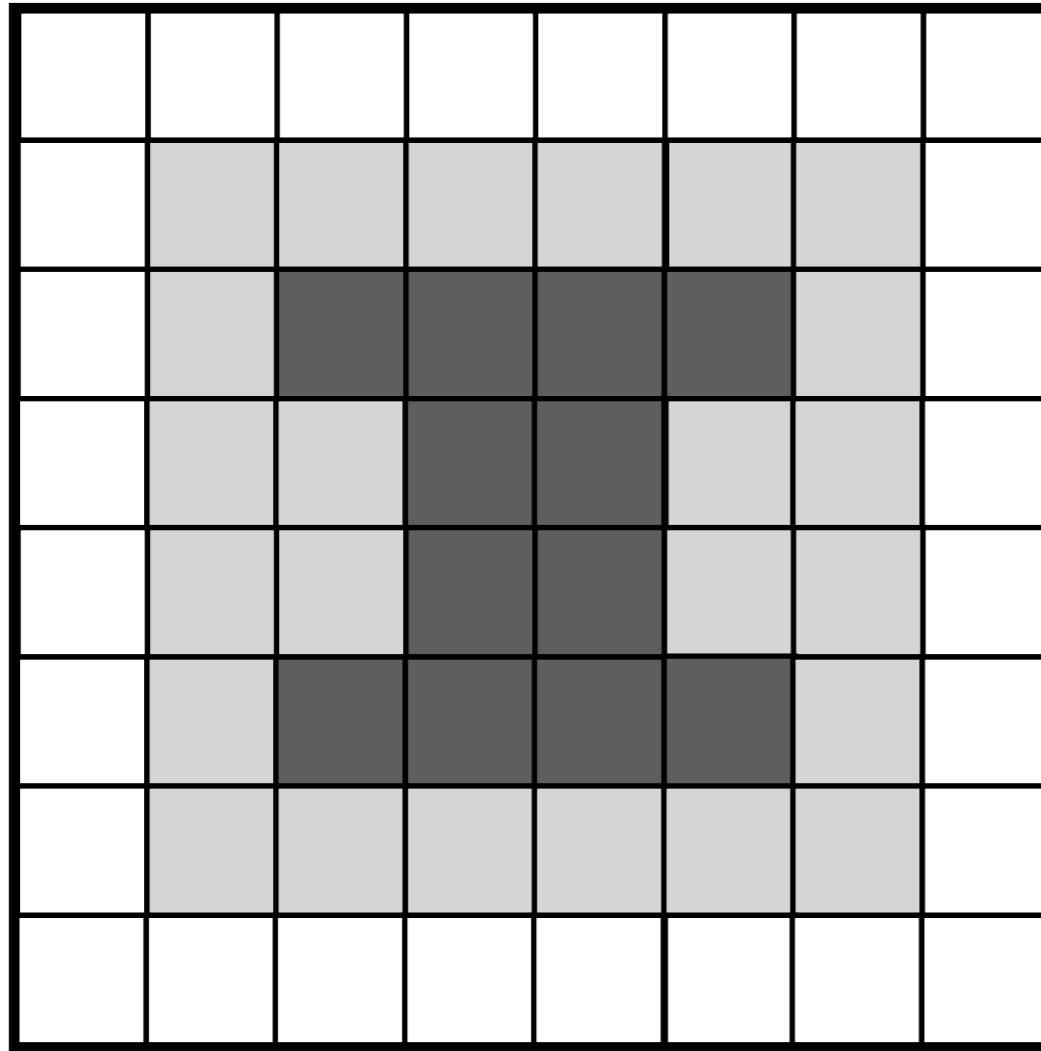




PRACTICAL MACHINE LEARNING: CLASSIFICATION & CONVOLUTIONAL NEURAL NETWORKS

Dr. Alexander Booth (he / him)
Universidad Católica del Norte, Chile
October 10th, 2024

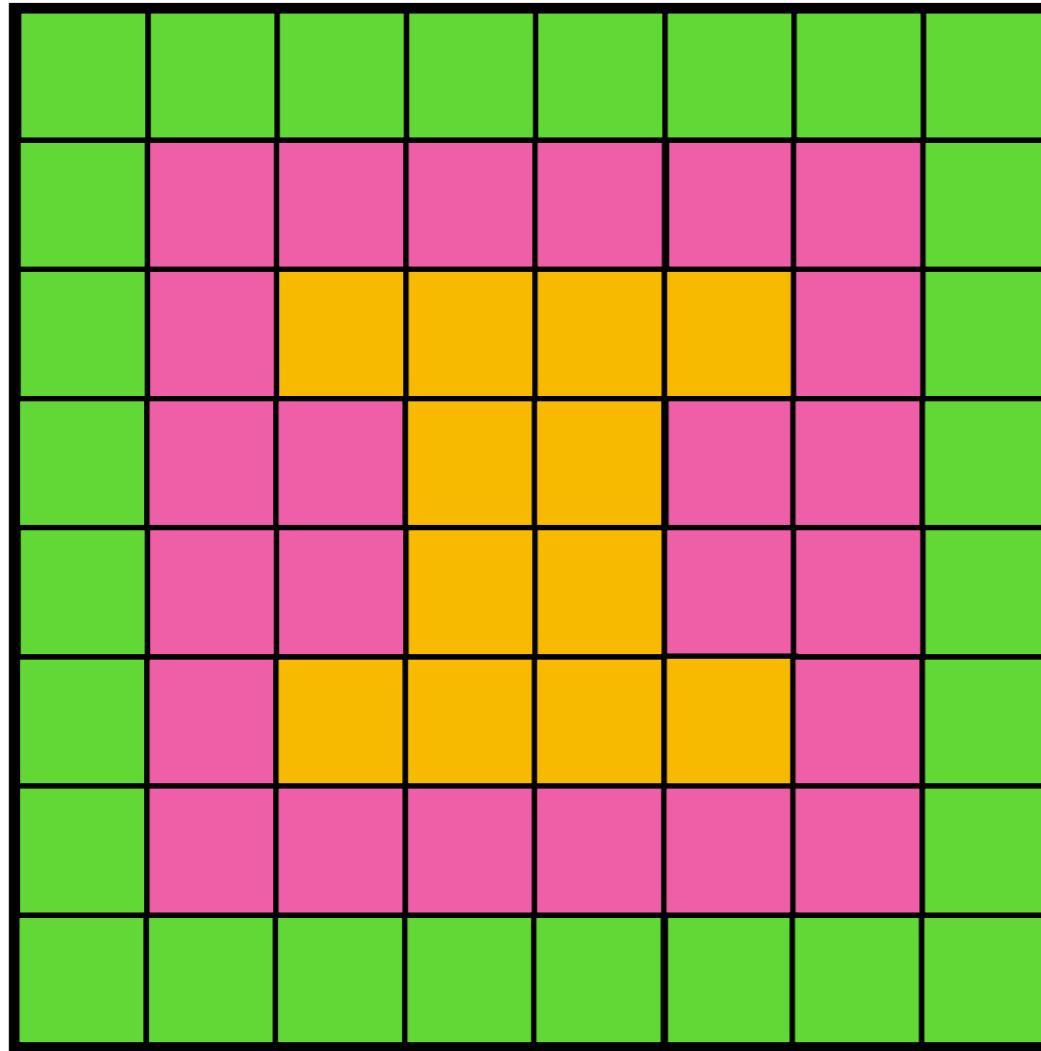
Fundamentals of Images



| | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

- Images are **grids of data**.
- Each cell in the grid is a **pixel**.
- The intensity of each pixel in a **greyscale** image can be quantified by a **single number** (8×8 image -> quantified by 64 numbers).

Fundamentals of Images

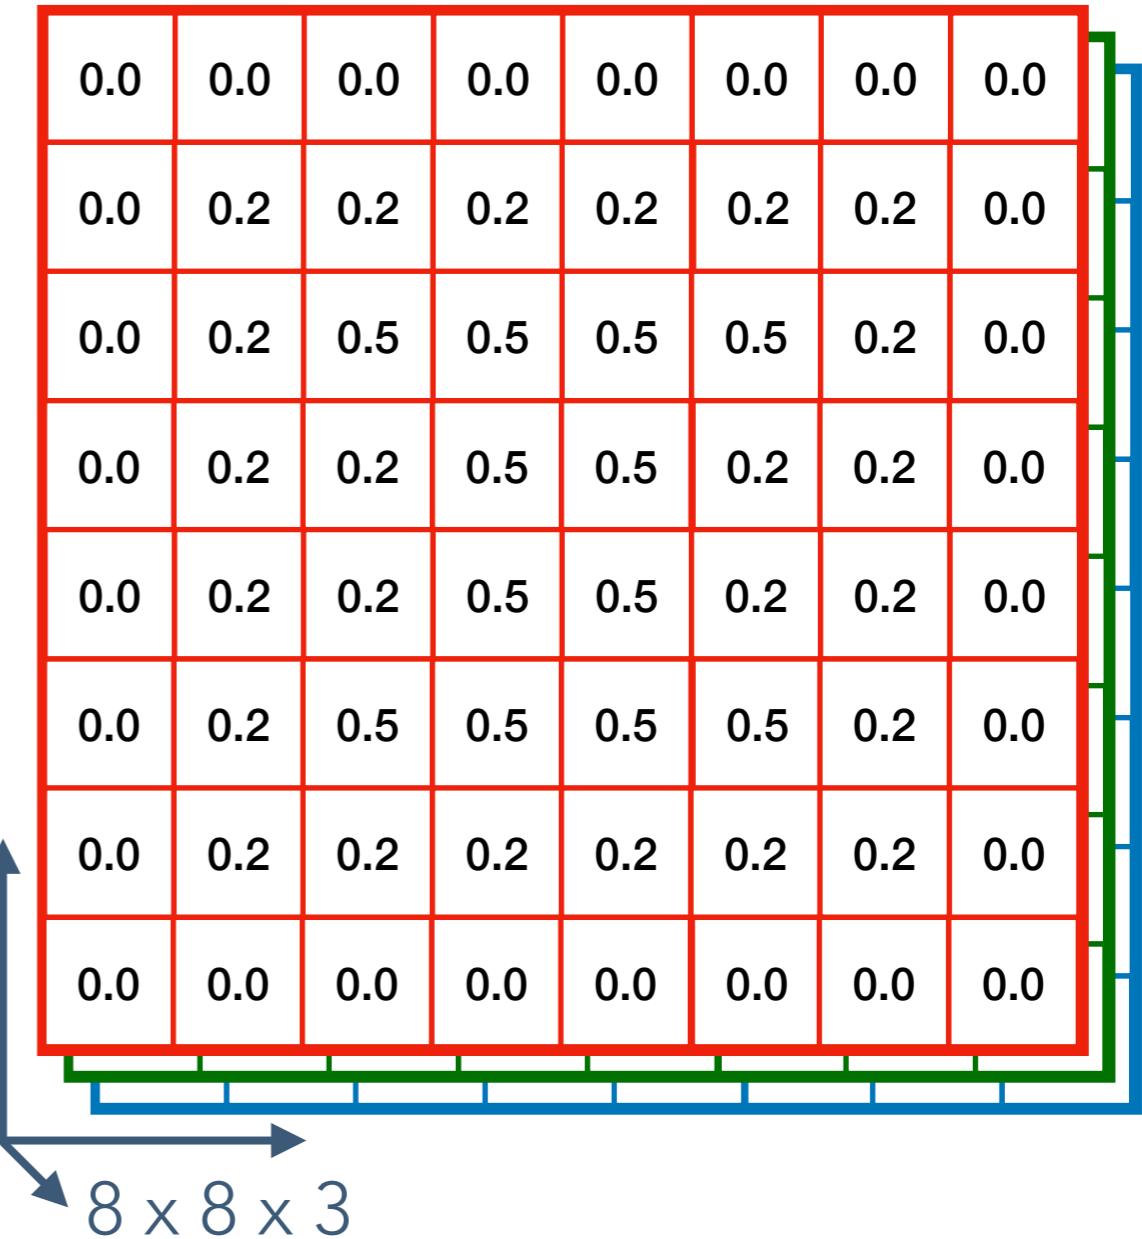
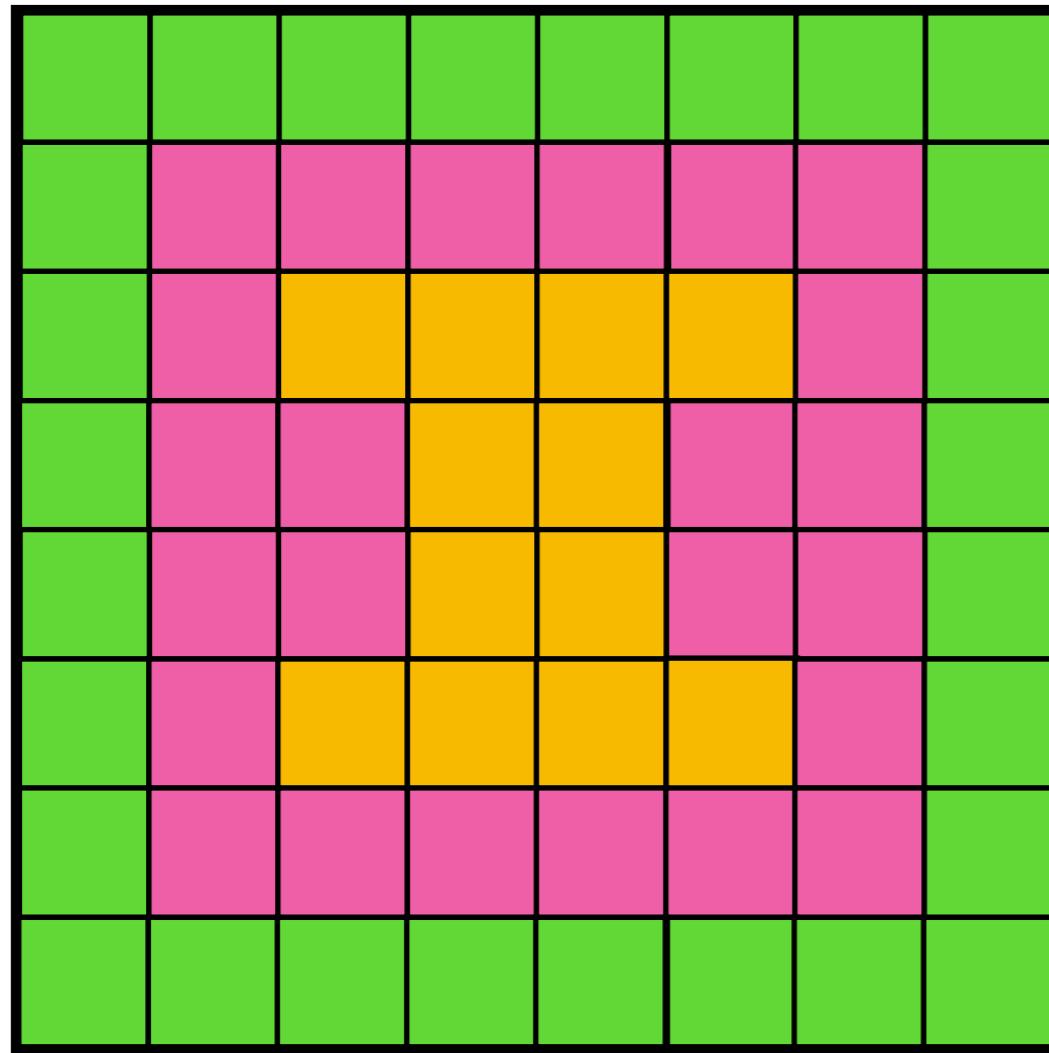


| | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |

8 x 8 x 3

- Images are **grids of data**.
- Each cell in the grid is a **pixel**.
- The intensity of each pixel in a **colour** image can be quantified by **three numbers, RGB** ($8 \times 8 \times 3$ image -> quantified by 192 numbers).

Computer Vision

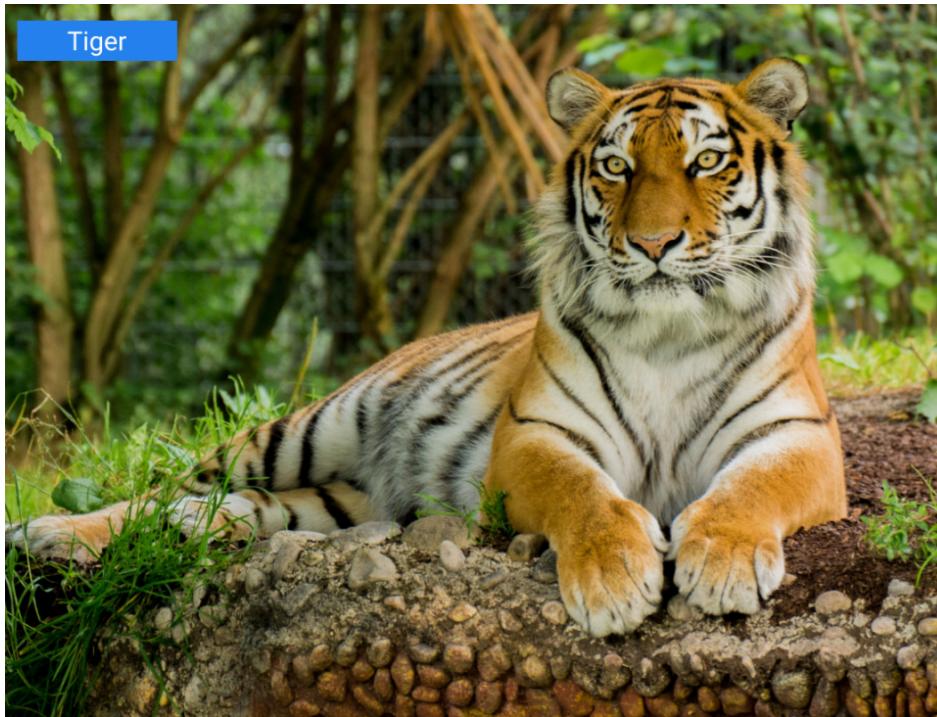


- Neural networks can be used to tackle many image related tasks -> “computer vision”:
 - ▶ Classification.
 - ▶ Object detection.
 - ▶ Semantic segmentation.
 - ▶ Instance segmentation.

Computer Vision



Examples & images taken from [1]



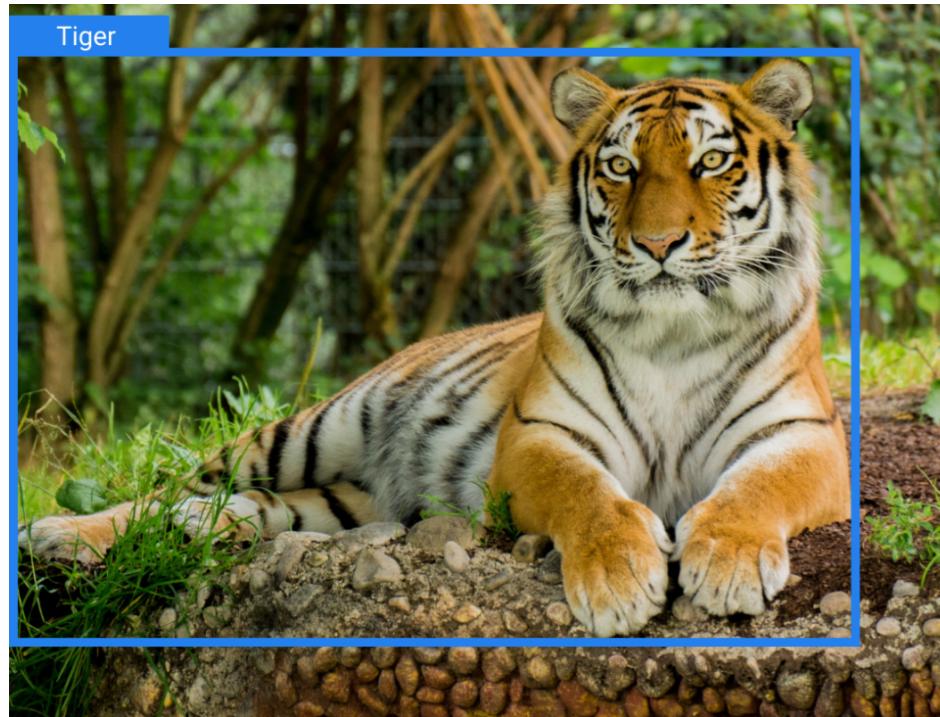
Categorising **entire** images into user-defined classes.

- Neural networks can be used to tackle many image related tasks -> “computer vision”:
 - ▶ **Classification.**
 - ▶ Object detection.
 - ▶ Semantic segmentation.
 - ▶ Instance segmentation.

Computer Vision



Examples & images taken from [1]



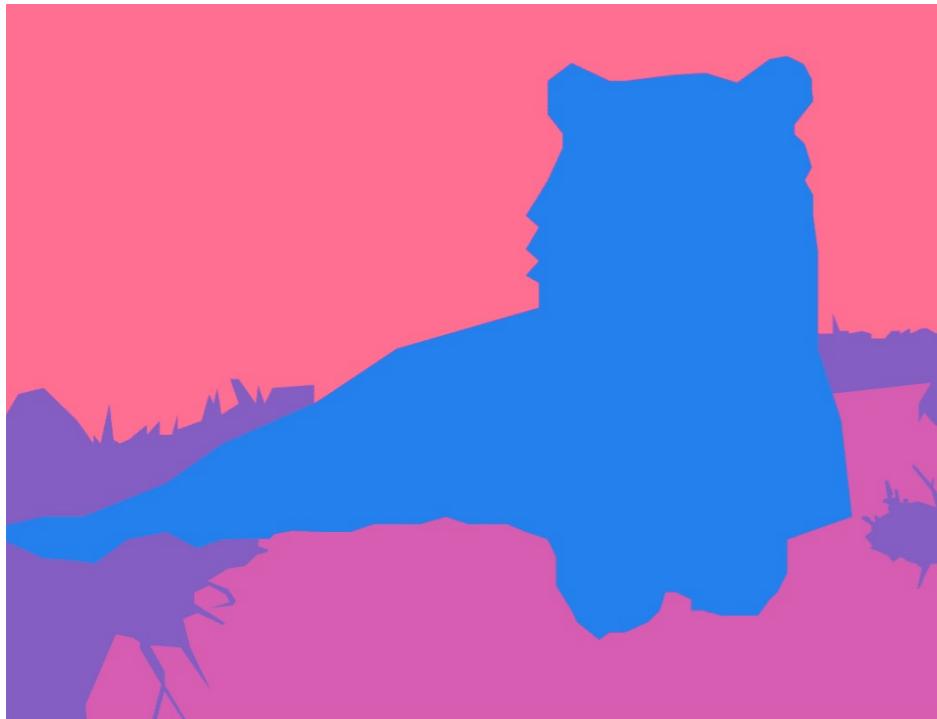
Identifying specific objects within an image, can be multiple objects (usually with bounding boxes).

- Neural networks can be used to tackle many image related tasks -> “computer vision”:
 - ▶ Classification.
 - ▶ **Object detection.**
 - ▶ Semantic segmentation.
 - ▶ Instance segmentation.

Computer Vision



Examples & images taken from [1]



Identifying each pixel in an image for more detailed classification.

- Neural networks can be used to tackle many image related tasks -> “computer vision”:
 - ▶ Classification.
 - ▶ Object detection.
 - ▶ **Semantic segmentation.**
 - ▶ Instance segmentation.



Computer Vision

Examples & images taken from [1]



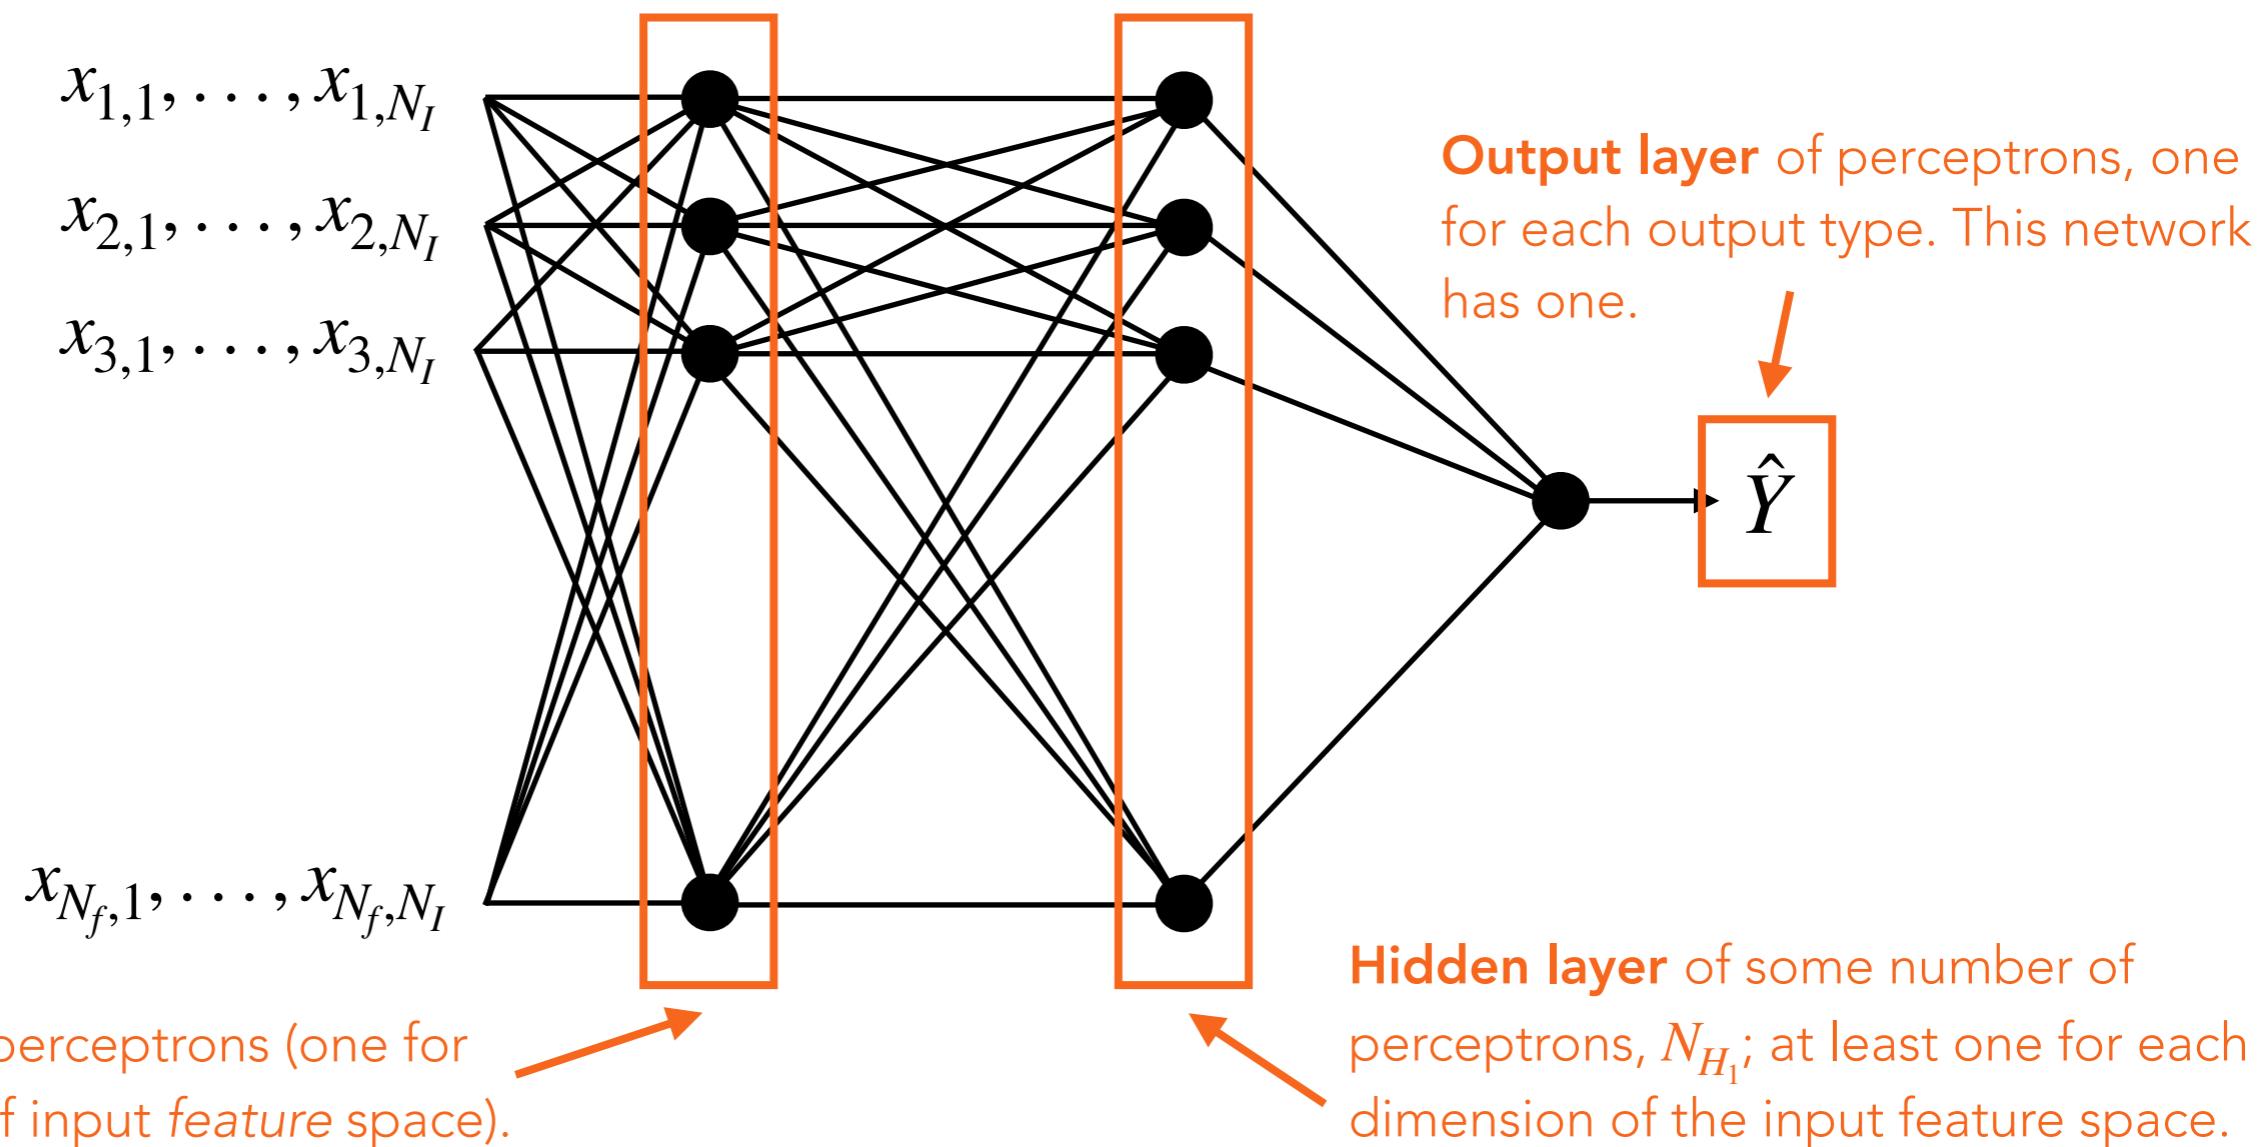
Recognise different instances of the same class.

- Neural networks can be used to tackle many image related tasks -> “computer vision”:
 - ▶ Classification.
 - ▶ Object detection.
 - ▶ Semantic segmentation.
 - ▶ **Instance segmentation.**

Recap: Basic Structure of a Neural Network



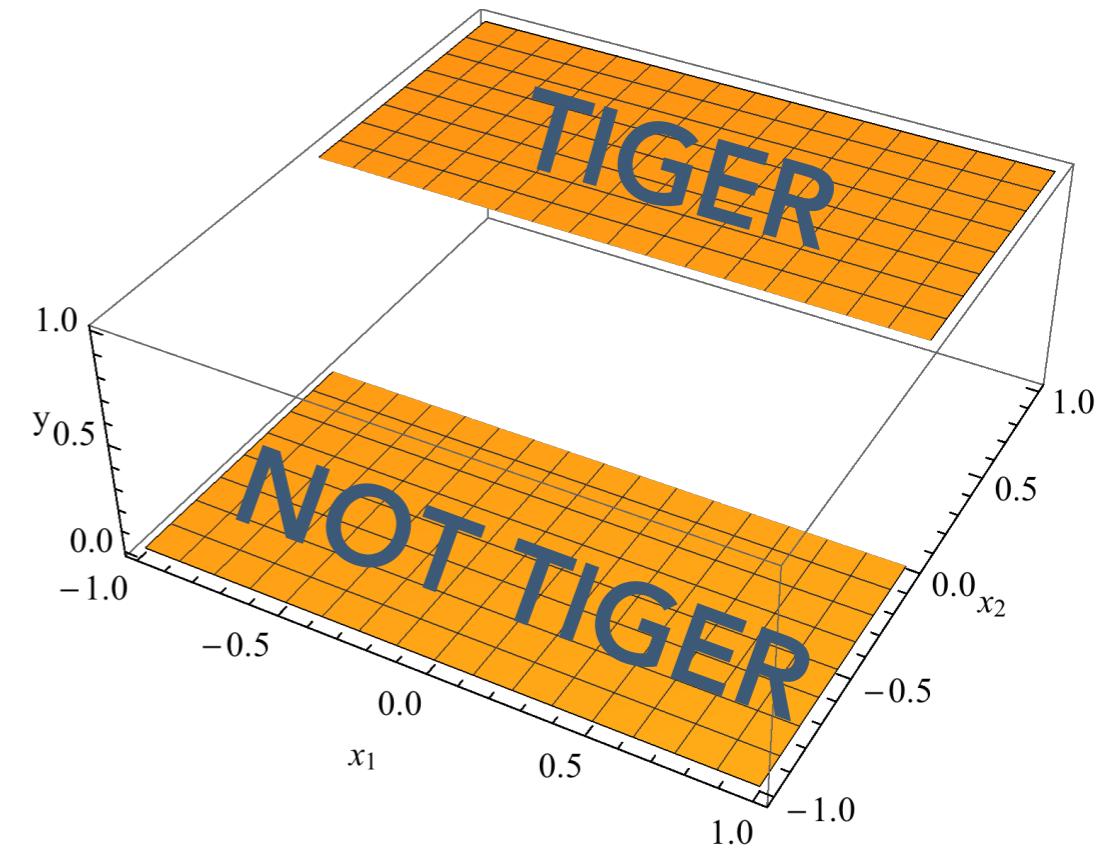
Initial features $\longleftrightarrow X = \begin{bmatrix} x_{1,1} & x_{1,2} & x_{1,3} & \dots & x_{1,N_I} \\ x_{2,1} & x_{2,2} & x_{2,3} & \dots & x_{2,N_I} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ x_{N_f,1} & x_{N_f,2} & x_{N_f,3} & \dots & x_{N_f,N_I} \end{bmatrix}$ Instances \downarrow Target feature (to predict) $Y = [y_{I_1}, y_{I_2}, \dots, y_{I_{N_I}}]$



Recap: Basic Structure of a Neural Network



- The final perceptron can be used to assign a type to data instance -> “classify” it.
- For example, consider a final perceptron with a binary activation function** that gives an all or nothing response with data that contains **two classes** of event (tiger and not tiger).
 - ▶ Let “all” response correspond to one type: tiger.
 - ▶ Let “nothing” response correspond to the other: not tiger



**Slightly more subtle -> usually use sigmoid.



Multi-class Output

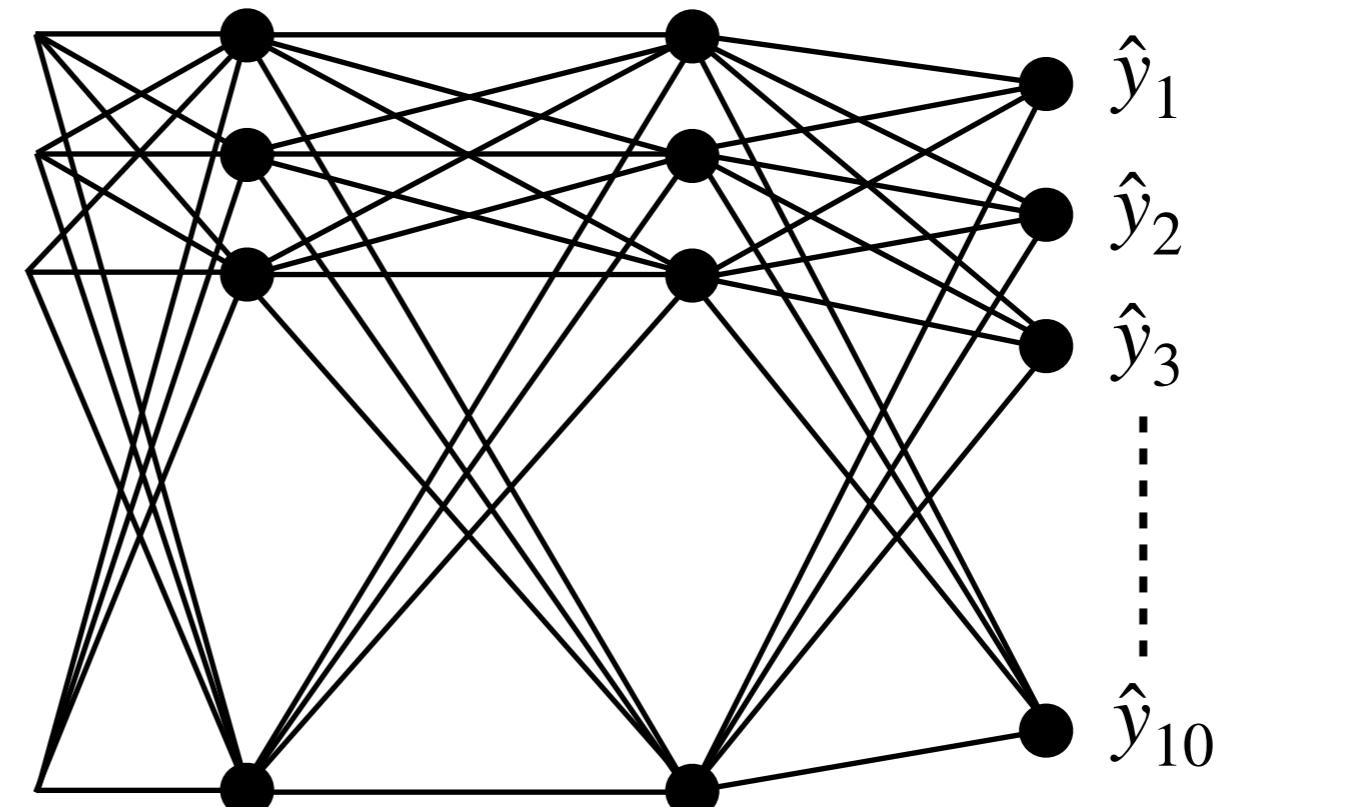
- Often have more than two possible classes - require a network with **an output layer with number of nodes equal to the number of possible classes.**
- For example: classifying the **MNIST dataset** [2]:
 - ▶ Collection of labelled images each containing a single handwritten digit between 0 and 9.
 - ▶ 70000 instances.
 - ▶ 28 x 28 greyscale pixels.
 - ▶ Pixel intensity in range [0, 255].

5 0 4 1 9 2 1 3 1 4
3 5 3 6 1 7 2 8 6 9
4 0 9 1 1 2 4 3 2 7
3 8 6 9 0 5 6 0 7 6
1 8 7 9 3 9 8 5 9 3
3 0 7 4 9 8 0 9 4 1
4 4 6 0 4 5 6 1 0 0
1 7 1 6 3 0 2 1 1 7
9 0 2 6 7 8 3 9 0 4
6 7 4 6 8 0 7 8 3 1 [3]



Multi-class Output

- Often have more than two possible classes - require a network with **an output layer with number of nodes equal to the number of possible classes.**
- For example: classifying the **MNIST dataset** [2]:
 - ▶ Collection of labelled images each containing a single handwritten digit between 0 and 9. ←
 - ▶ 70000 instances.
 - ▶ 28×28 greyscale pixels.
 - ▶ Pixel intensity in range [0, 255].

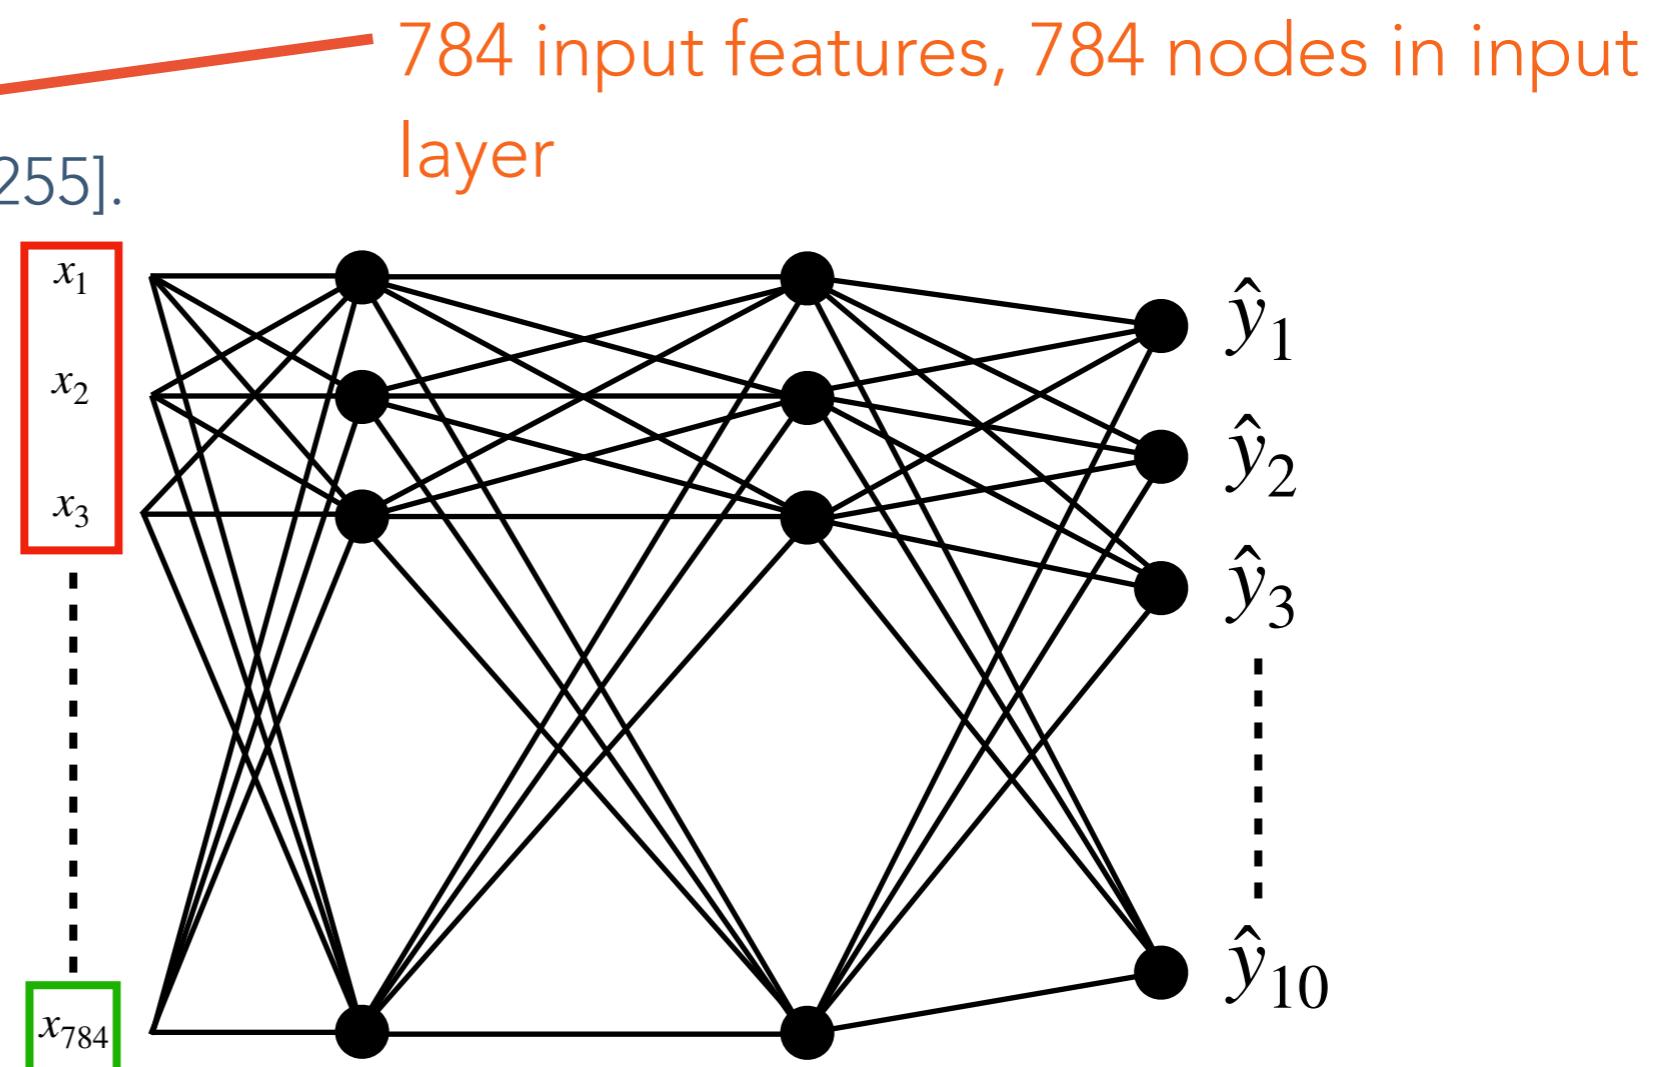




Multi-class Output

- Often have more than two possible classes - require a network with **an output layer with number of nodes equal to the number of possible classes.**
- For example: classifying the **MNIST dataset** [2]:
 - Collection of labelled images each containing a single handwritten digit between 0 and 9.
 - 70000 instances.
 - 28 x 28 greyscale pixels.
 - Pixel intensity in range [0, 255].

| | | | | | | | |
|-----------|-----------|-----|-----|-----|-----|-----------|-----------|
| x_1 | x_2 | ... | ... | ... | ... | x_{27} | x_{28} |
| x_{29} | ... | ... | ... | ... | ... | x_{56} | |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| x_{729} | ... | ... | ... | ... | ... | x_{756} | |
| x_{757} | x_{758} | ... | ... | ... | ... | x_{783} | x_{784} |

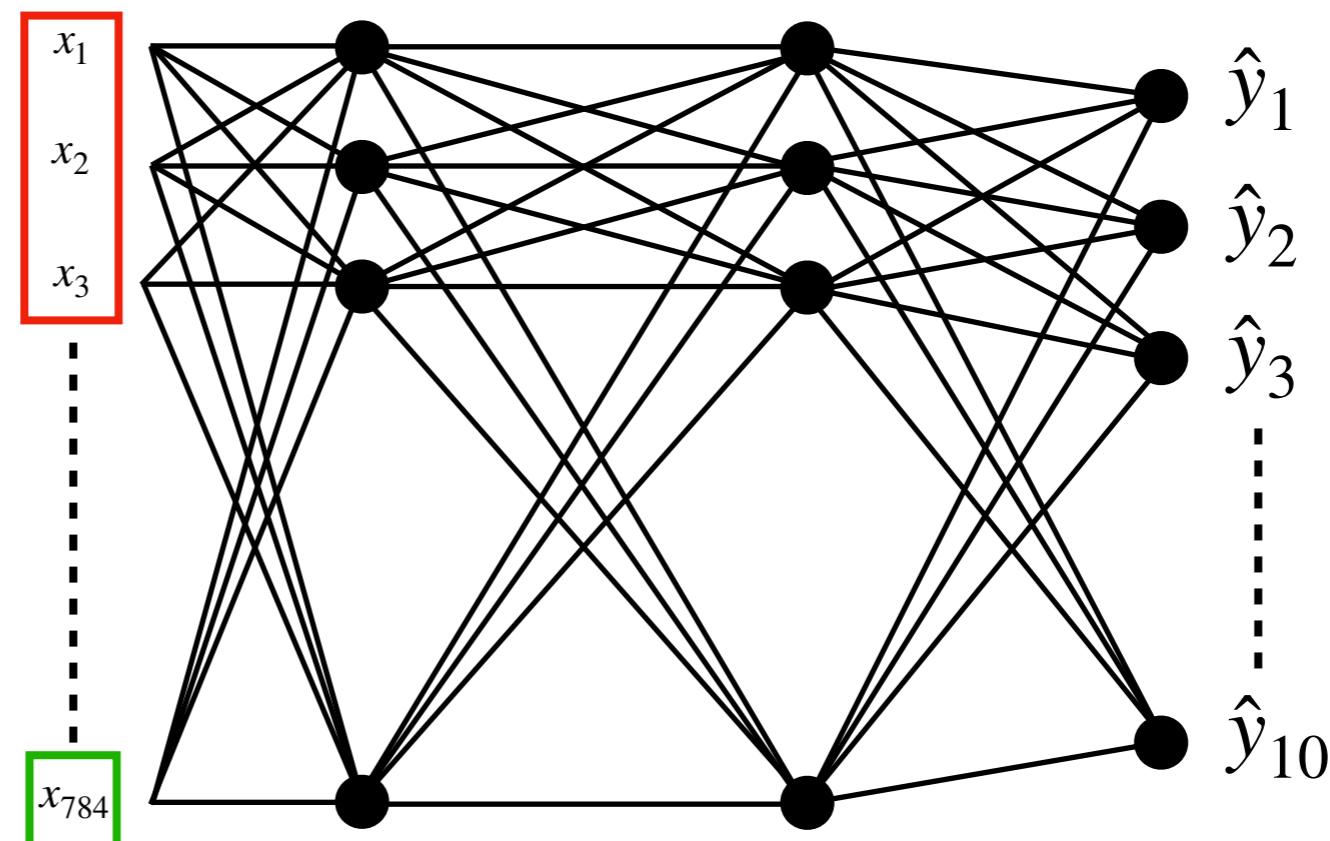




Multi-class Output

- From each output node:
 - \hat{y}_1 makes the decision if an instance is the number 1 vs is not the number 1.
 - \hat{y}_2 makes the decision if an instance is the number 2 vs is not the number 2.
 - etc...
- Called **one vs. all** model.
 - In this approach each binary set of outcomes is determined independently of all other sets.
 - E.g. determining outcome of 1 vs. not 1 without considering 2 vs. not 2.

| | | | | | | | |
|-----------|-----------|-----|-----|-----|-----|-----------|-----------|
| x_1 | x_2 | ... | ... | ... | ... | x_{27} | x_{28} |
| x_{29} | ... | ... | ... | ... | ... | x_{56} | |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| x_{729} | ... | ... | ... | ... | ... | x_{756} | |
| x_{757} | x_{758} | ... | ... | ... | ... | x_{783} | x_{784} |





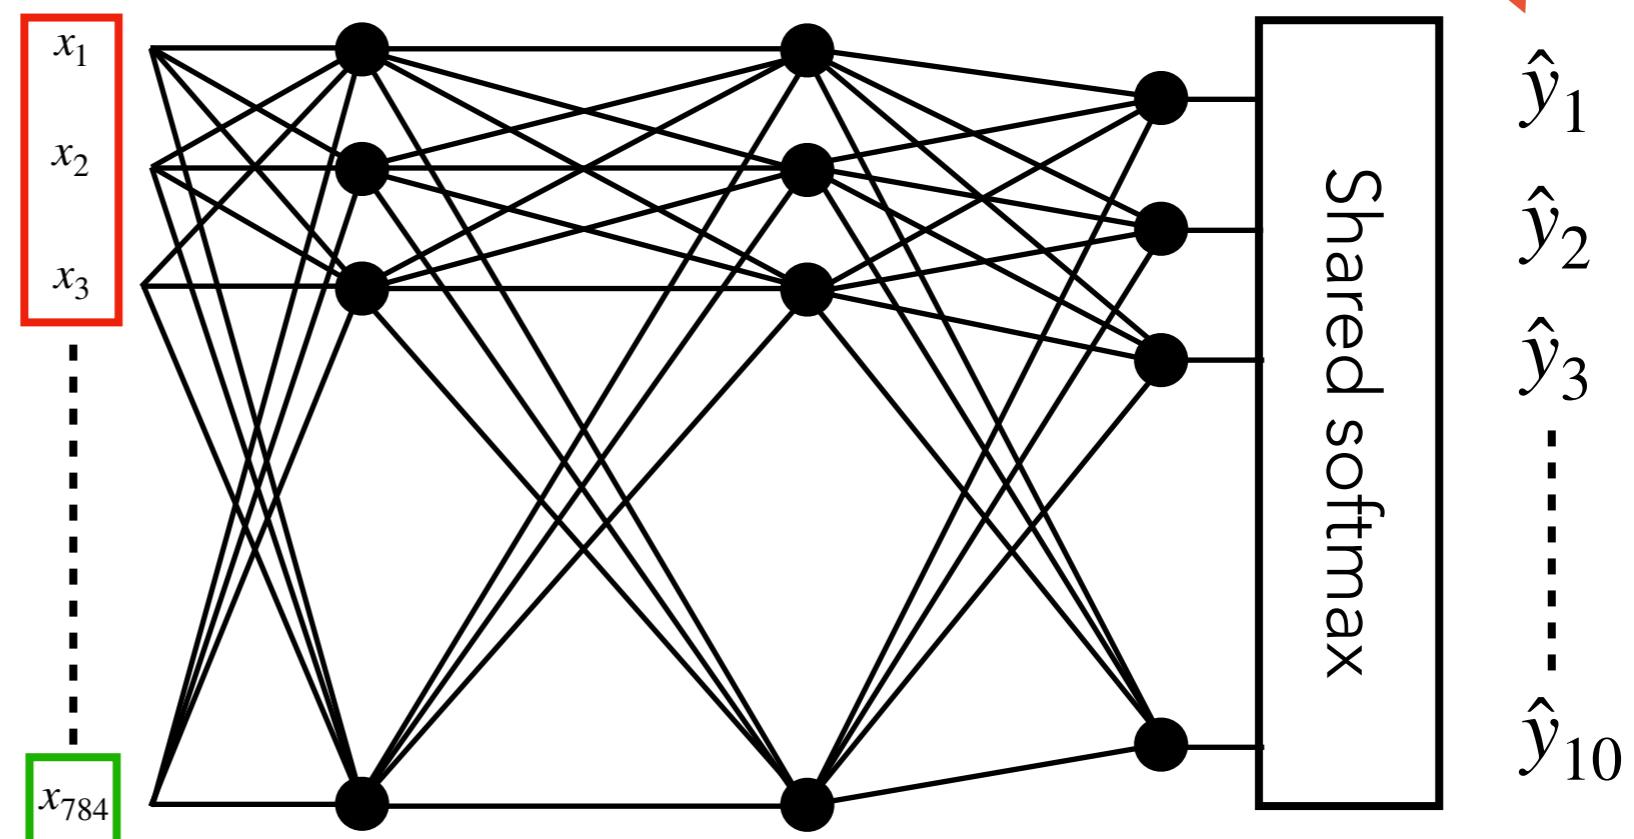
Multi-class Output

- Usually want to predict the probabilities of each outcome relative to each other -> **one vs. one** model.
 - In this approach after the output layer we can apply the **softmax function** over all outputs with respect to each individual output.
 - Transforms / renormalises the output
 - If j labels the class and there are N_C classes:

$$\hat{y}_j(x) = \frac{e^{W_j^T x}}{\sum_k^{N_C} e^{w_k^T x}}$$

$$\sum_j^{N_C} \hat{y}_j = 1$$

| | | | | | | | |
|-----------|-----------|---------|---------|---------|---------|-----------|-----------|
| x_1 | x_2 | \dots | \dots | \dots | \dots | x_{27} | x_{28} |
| x_{29} | \dots | \dots | \dots | \dots | \dots | x_{56} | |
| \dots | \dots | \dots | \dots | \dots | \dots | \dots | |
| \dots | \dots | \dots | \dots | \dots | \dots | \dots | |
| \dots | \dots | \dots | \dots | \dots | \dots | \dots | |
| \dots | \dots | \dots | \dots | \dots | \dots | \dots | |
| x_{729} | \dots | \dots | \dots | \dots | \dots | \dots | x_{756} |
| x_{757} | x_{758} | \dots | \dots | \dots | \dots | x_{783} | x_{784} |





Performance on MNIST

- The MNIST website [2] lists a number of complicated ways to train neural networks to solve this classification problem.
- MLPs have produced good results! Best error rate achieved of 0.35% (Ciresan et al [4]).

| ID | architecture (number of neurons in each layer) | test error for | | best test error [%] | simulation time [h] | weights [milions] |
|----|---|---------------------|-------------|------------------------|------------------------|----------------------|
| | | best validation [%] | | | | |
| 1 | 1000, 500, 10 | | 0.49 | 0.44 | 23.4 | 1.34 |
| 2 | 1500, 1000, 500, 10 | | 0.46 | 0.40 | 44.2 | 3.26 |
| 3 | 2000, 1500, 1000, 500, 10 | | 0.41 | 0.39 | 66.7 | 6.69 |
| 4 | 2500, 2000, 1500, 1000, 500, 10 | | 0.35 | 0.32 | 114.5 | 12.11 |

Performance on MNIST



35 instances out of 10000 in the test set!

| | |
|------------------------------------|--|
| Correct label | |
| NUMBER | |
| Two most likely predictions (L, R) | |

| | | | | | | |
|------------|------------|------------|------------|------------|------------|------------|
| 1 2 1 7 | 1 1 7 1 | 9 8 9 8 | 9 9 5 9 | 9 9 7 9 | 2 5 3 5 | 8 8 2 3 |
| 4 9 4 9 | 5 5 3 5 | 9 4 9 7 | 4 9 4 9 | 4 4 9 4 | 2 2 0 2 | 5 5 3 5 |
| 1 6 1 6 | 9 4 9 4 | 0 0 6 0 | 6 6 0 6 | 6 6 8 6 | 1 1 7 9 | 1 1 7 1 |
| 9 9 4 9 | 0 0 5 0 | 5 5 3 5 | 8 8 9 8 | 9 9 7 9 | 7 7 1 7 | 1 1 6 1 |
| 2 7 2 7 | 8 8 5 8 | 2 2 7 8 | 6 6 1 6 | 5 5 6 5 | 4 4 9 4 | 0 0 6 0 |

Convolutional Neural Networks (CNNs)

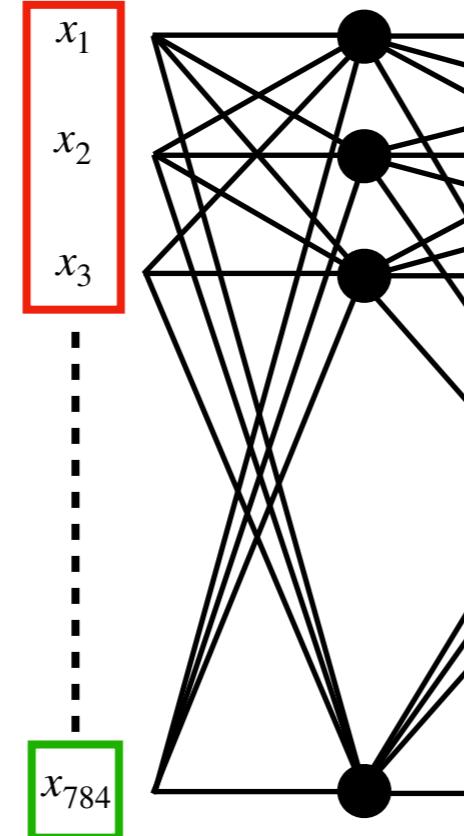


CNNs vs. MLPs

- Like MLPs, CNNs can be applied to problems in **computer vision**.
 - ▶ Particularly good at dealing with **grid-like data** (images).
- Unlike MLPs they can take advantage of **spacial correlations** in the input feature space.

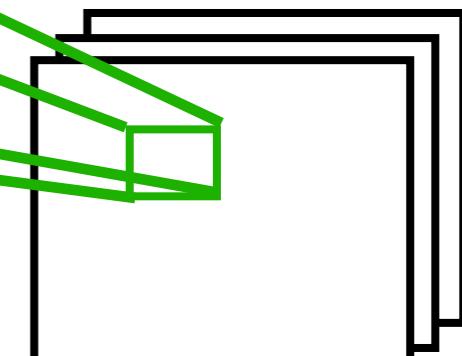
MLP

| | | | | | | | |
|-----------|-----------|-----|-----|-----|-----|-----------|-----------|
| x_1 | x_2 | ... | ... | ... | ... | x_{27} | x_{28} |
| x_{29} | ... | ... | ... | ... | ... | ... | x_{56} |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| x_{729} | ... | ... | ... | ... | ... | ... | x_{756} |
| x_{757} | x_{758} | ... | ... | ... | ... | x_{783} | x_{784} |



CNN

| | | | | | | | |
|-----------|-----------|-----|-----|-----|-----|-----------|-----------|
| x_1 | x_2 | ... | ... | ... | ... | x_{27} | x_{28} |
| x_{29} | ... | ... | ... | ... | ... | ... | x_{56} |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... |
| x_{729} | ... | ... | ... | ... | ... | ... | x_{756} |
| x_{757} | x_{758} | ... | ... | ... | ... | x_{783} | x_{784} |

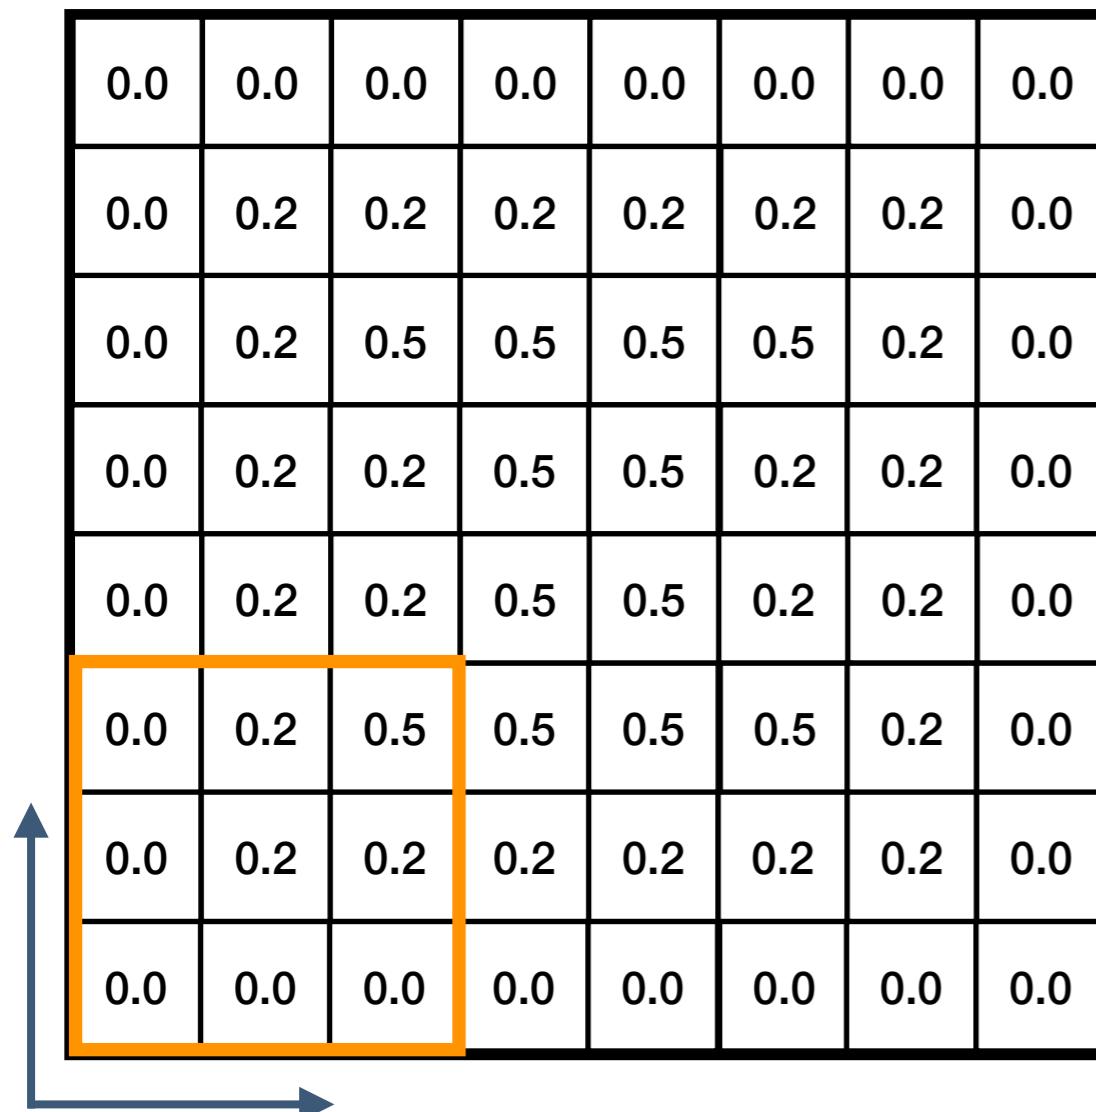


Convolutional layers



Convolution Layers

- In the convolutional layers we **perform convolutions** of the whole input image **piece by piece** by sliding **a filters** over the image.
- Elements of the **filters** are the **weights** in this network -> the thing that you optimise.



| | | |
|----------|----------|----------|
| w_{11} | w_{21} | w_{31} |
| w_{12} | w_{22} | w_{32} |
| w_{13} | w_{23} | w_{33} |

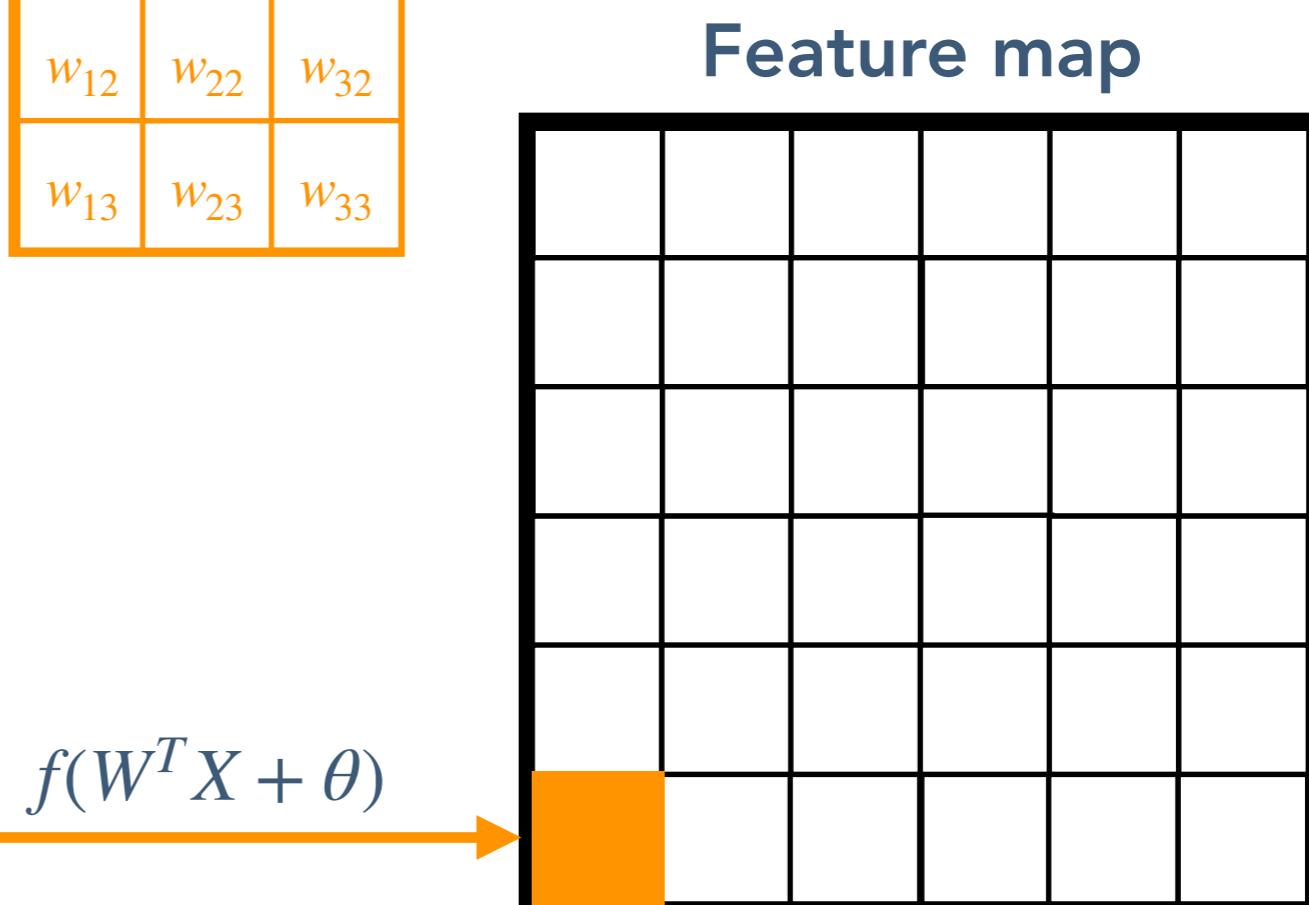
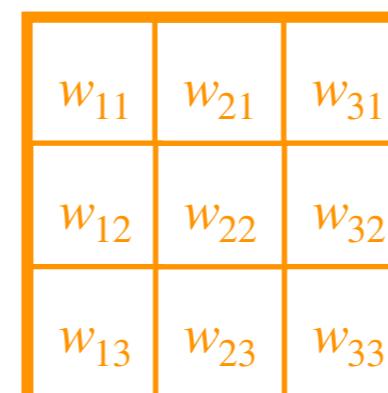
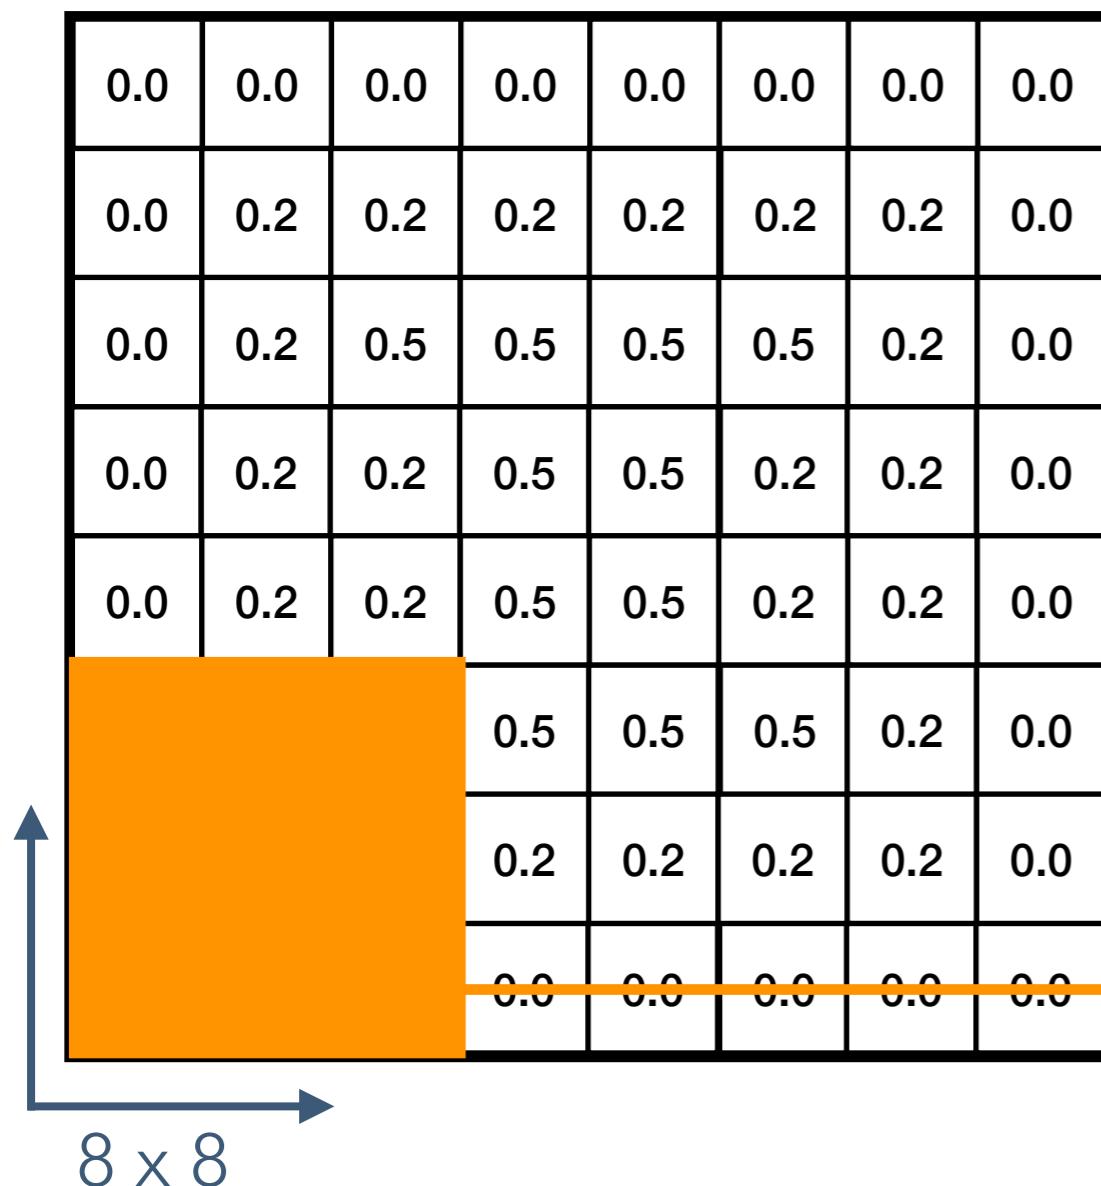
3 x 3 filter

Can be $N \times N$ for $N \leq M$



Convolution Layers

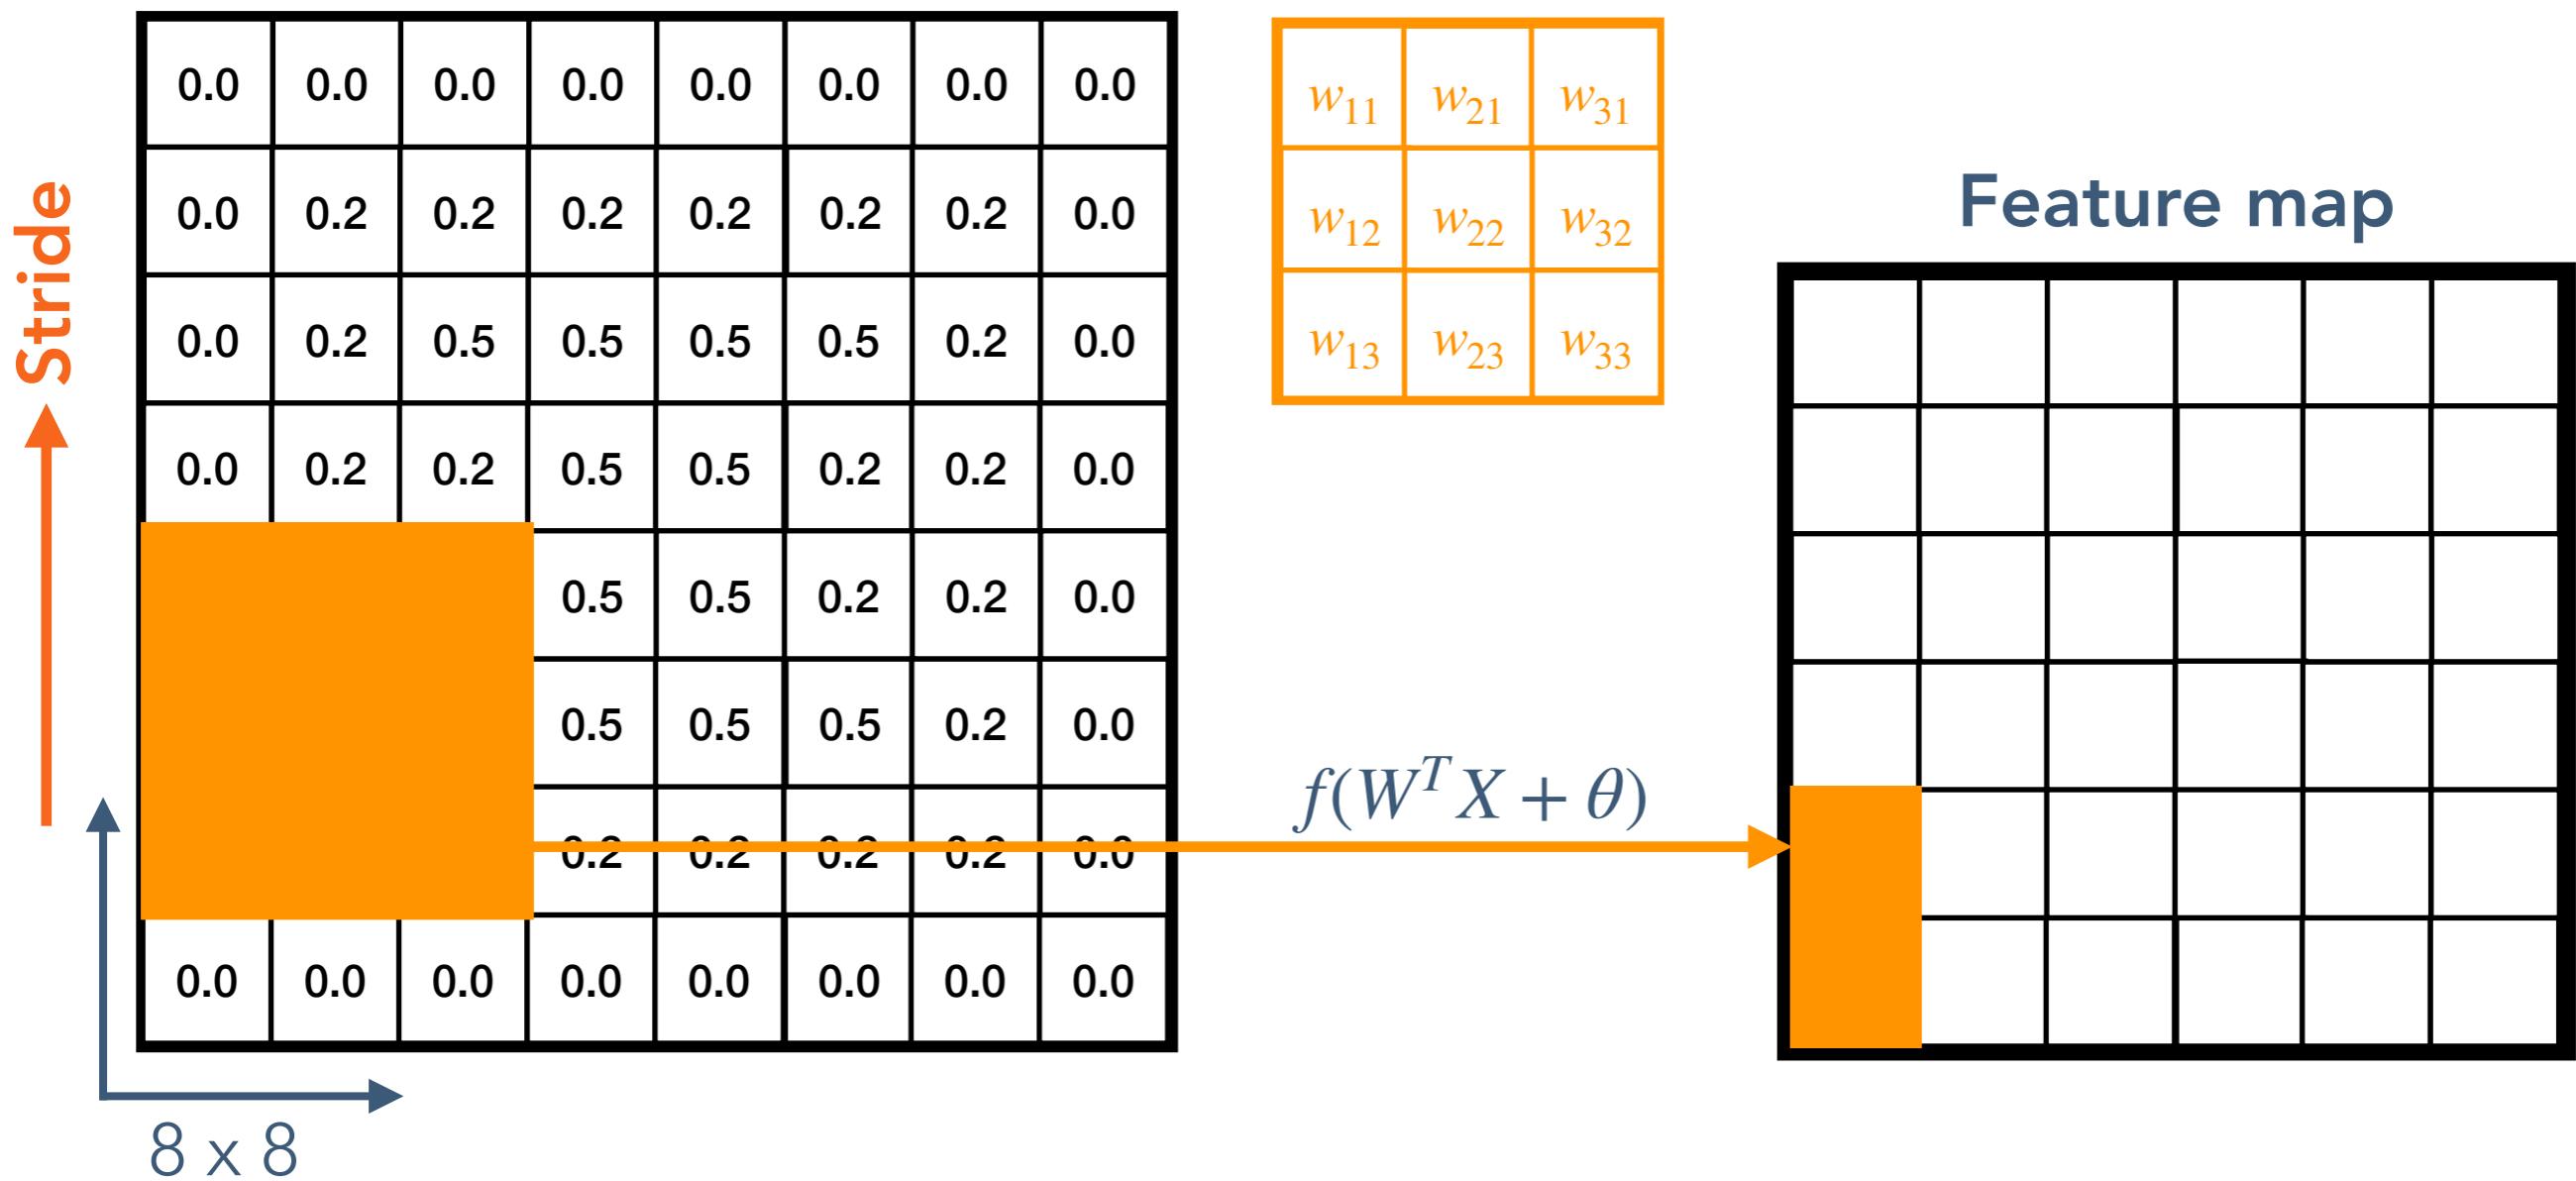
- In the convolutional layers we **perform convolutions** of the whole input image **piece by piece** by sliding **a filters** over the image.
- Apply the filter in “position 1” -> taking a weighted sum of pixels in the filter’s **receptive field**.





Convolution Layers

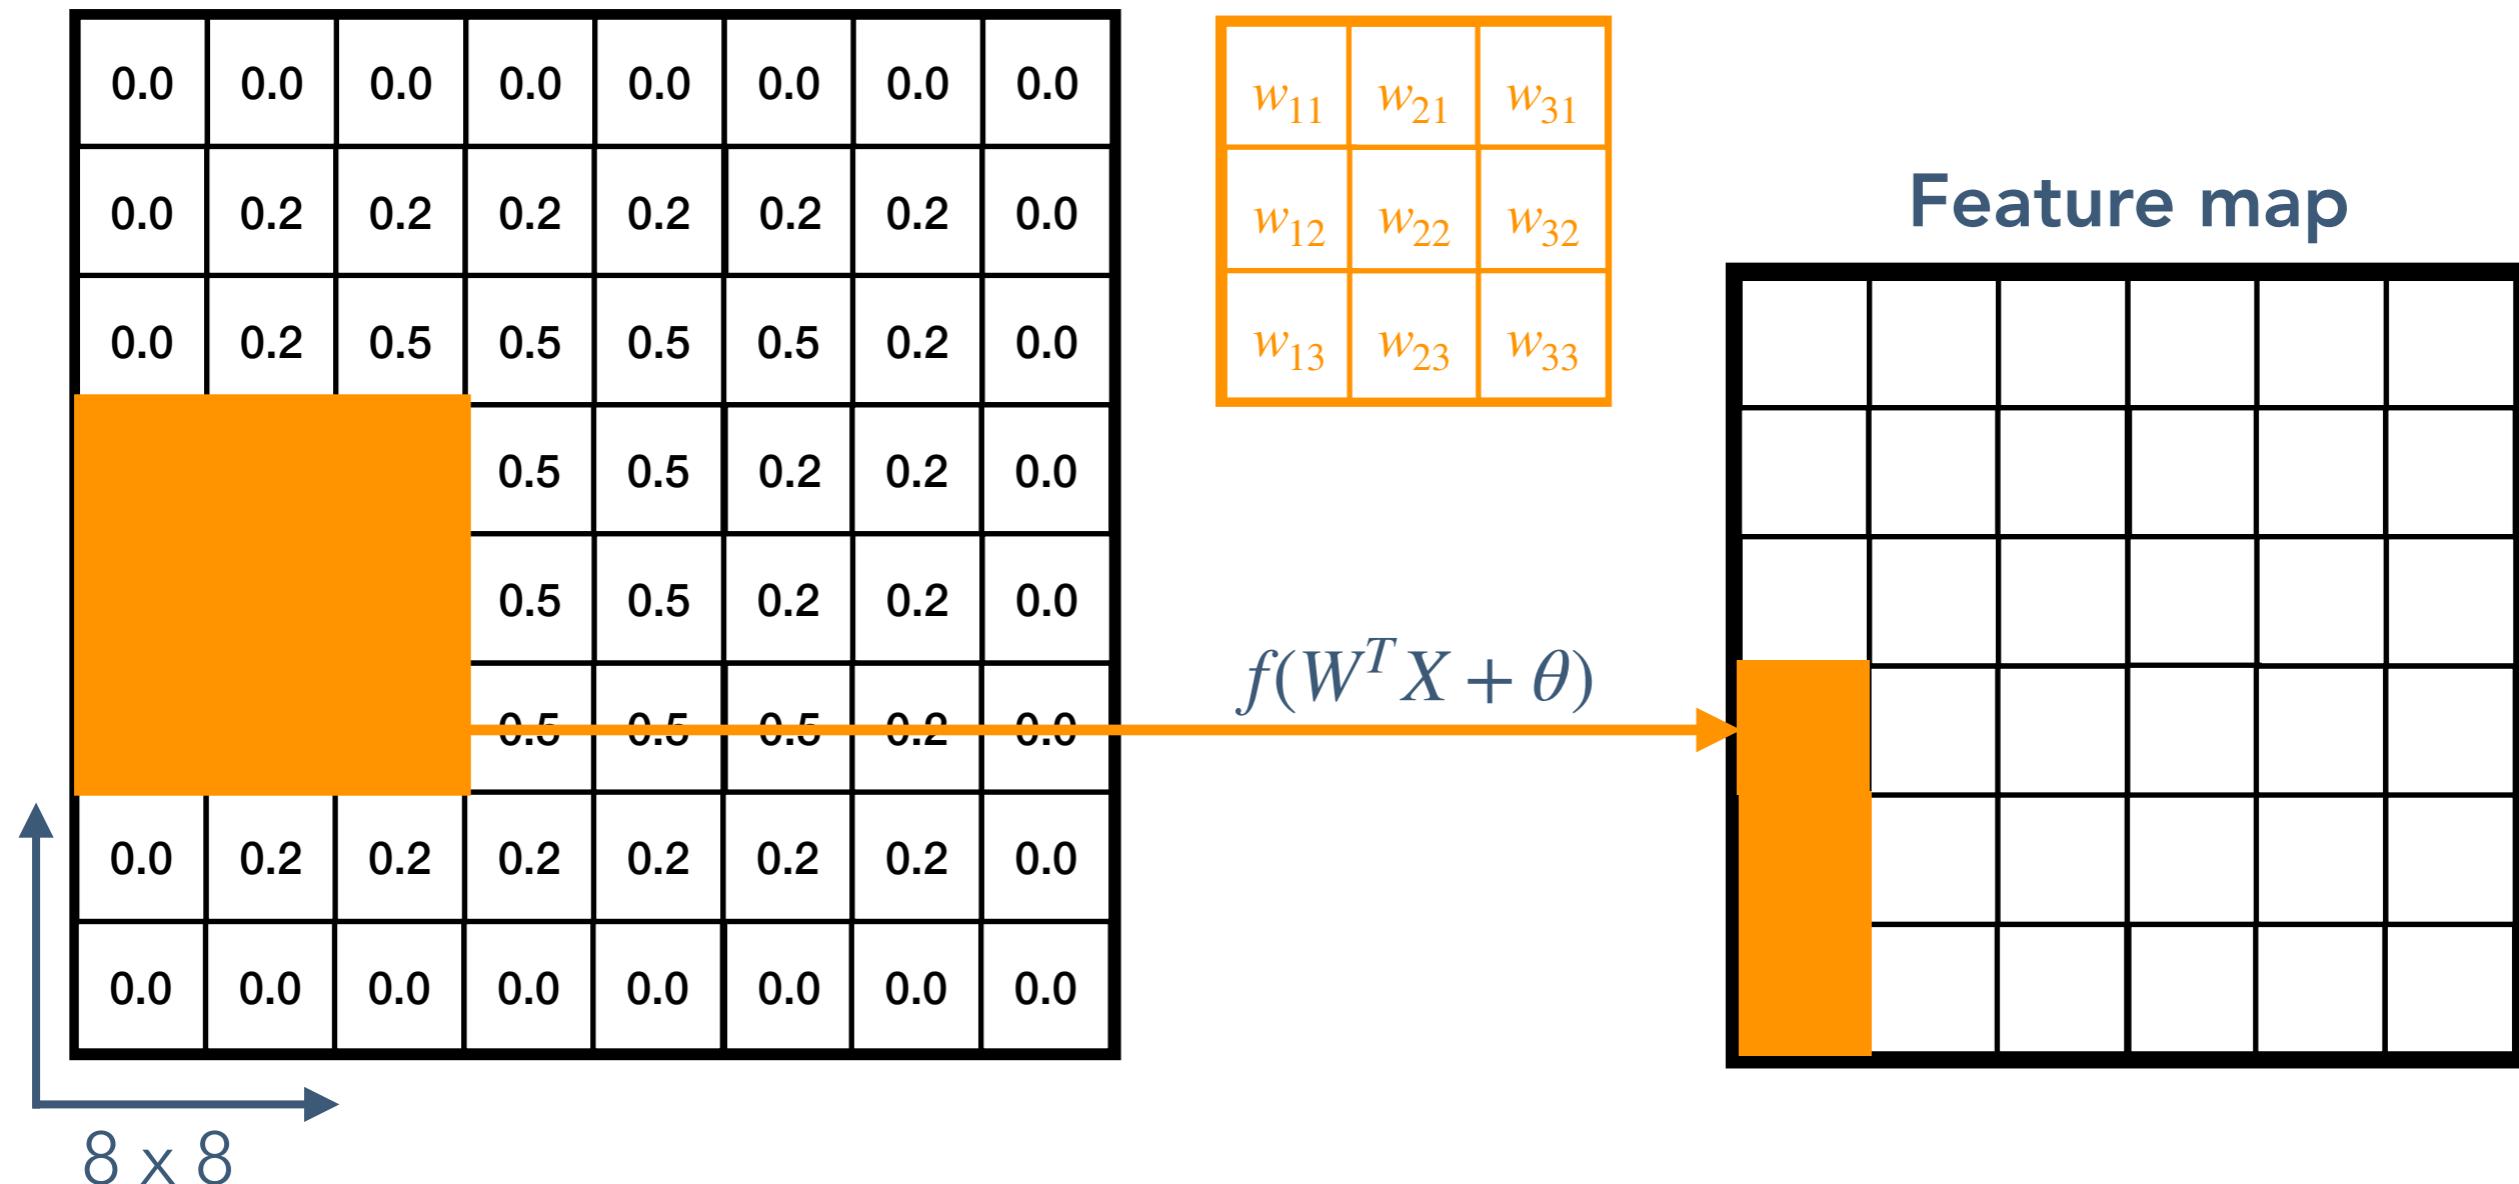
- In the convolutional layers we **perform convolutions** of the whole input image **piece by piece** by sliding **a filters** over the image.
- Apply the filter in position 2. Distance between position 1 and position 2 is governed by the **stride**.





Convolution Layers

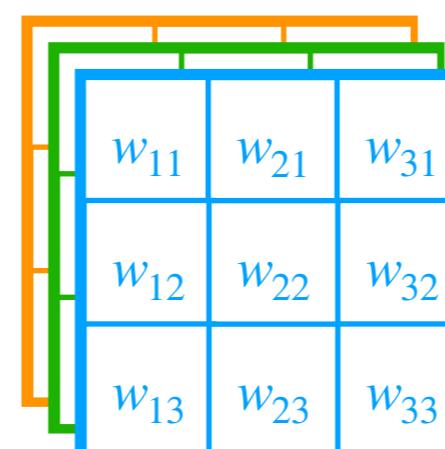
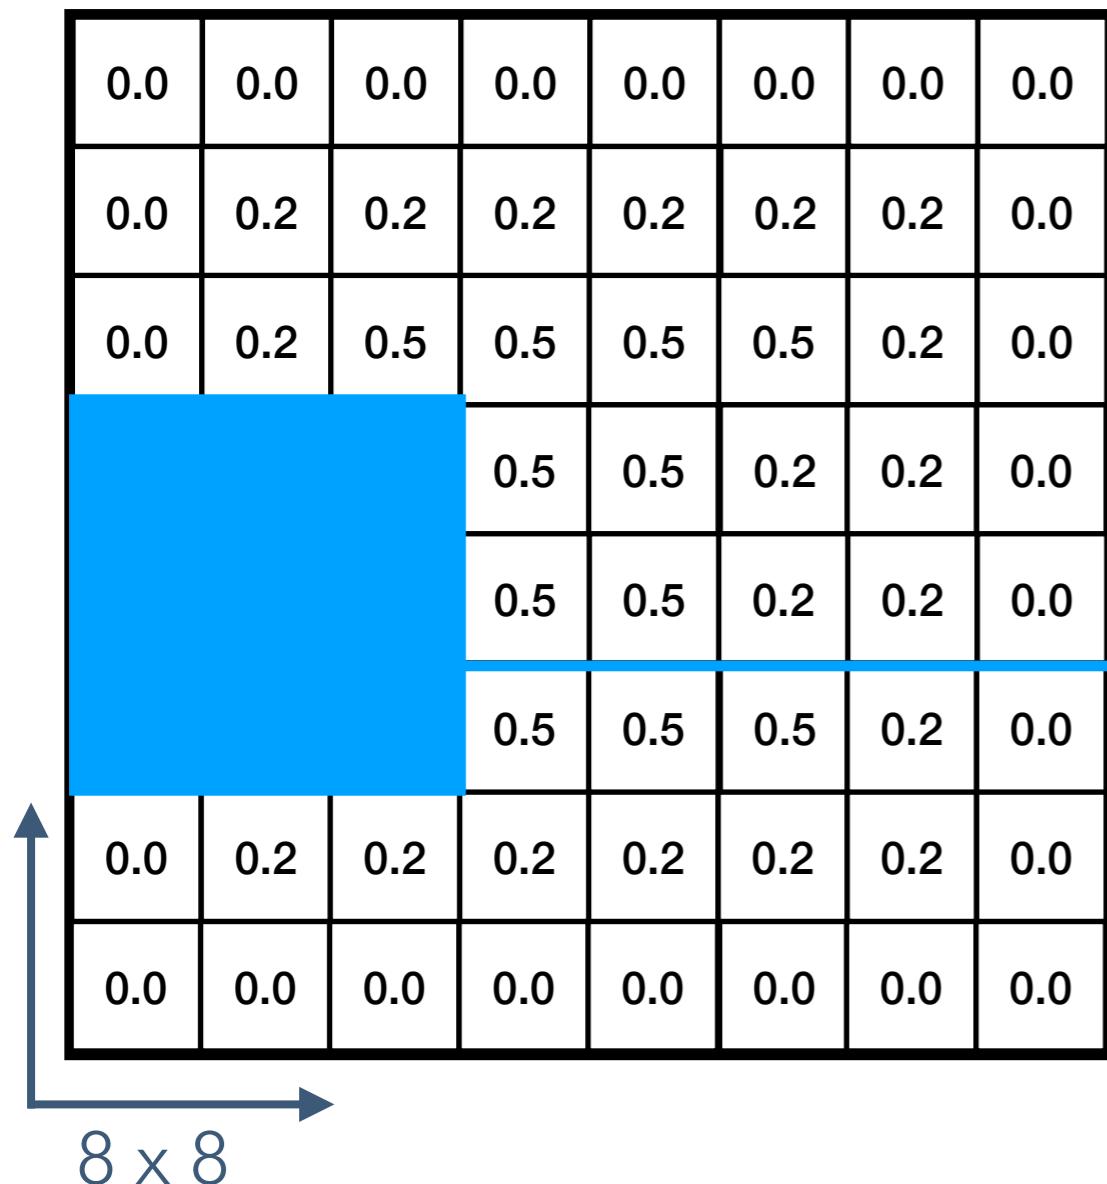
- Each “neuron” in a given feature map has the **same weights**.
 - ▶ Reduced numbers of parameters in the model compared to an MLP.
 - ▶ Once a CNN has learned to recognise a pattern in one location in the image, it can recognise it **anywhere** in the image - not true for an MLP.





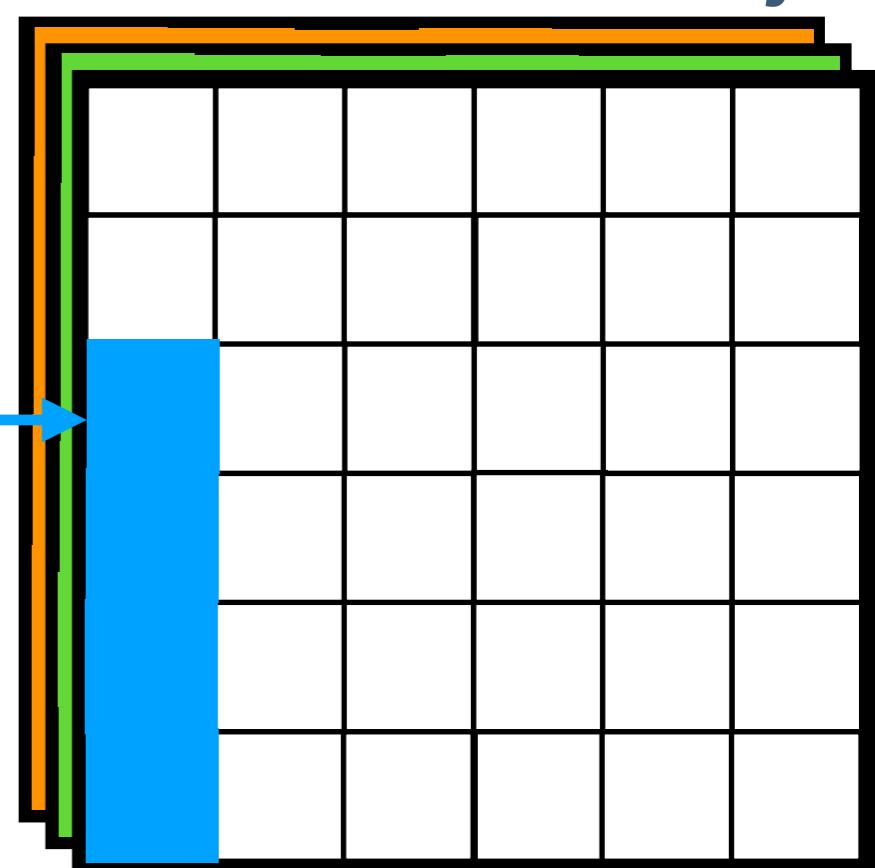
Convolution Layers

- A network architecture design choice is how many **feature maps** are produced per convolutional layer -> a different filter (and different set of weights) each.
- Each filter enhances a different feature in the image.



$$f(W^T X + \theta)$$

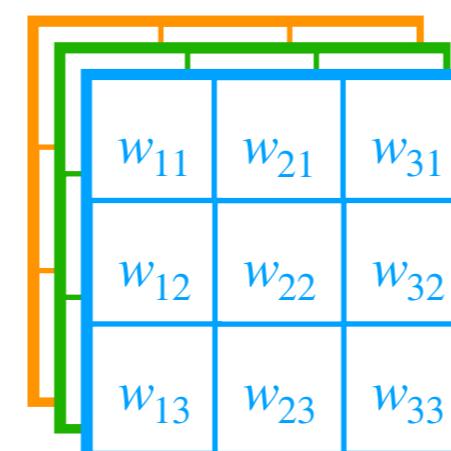
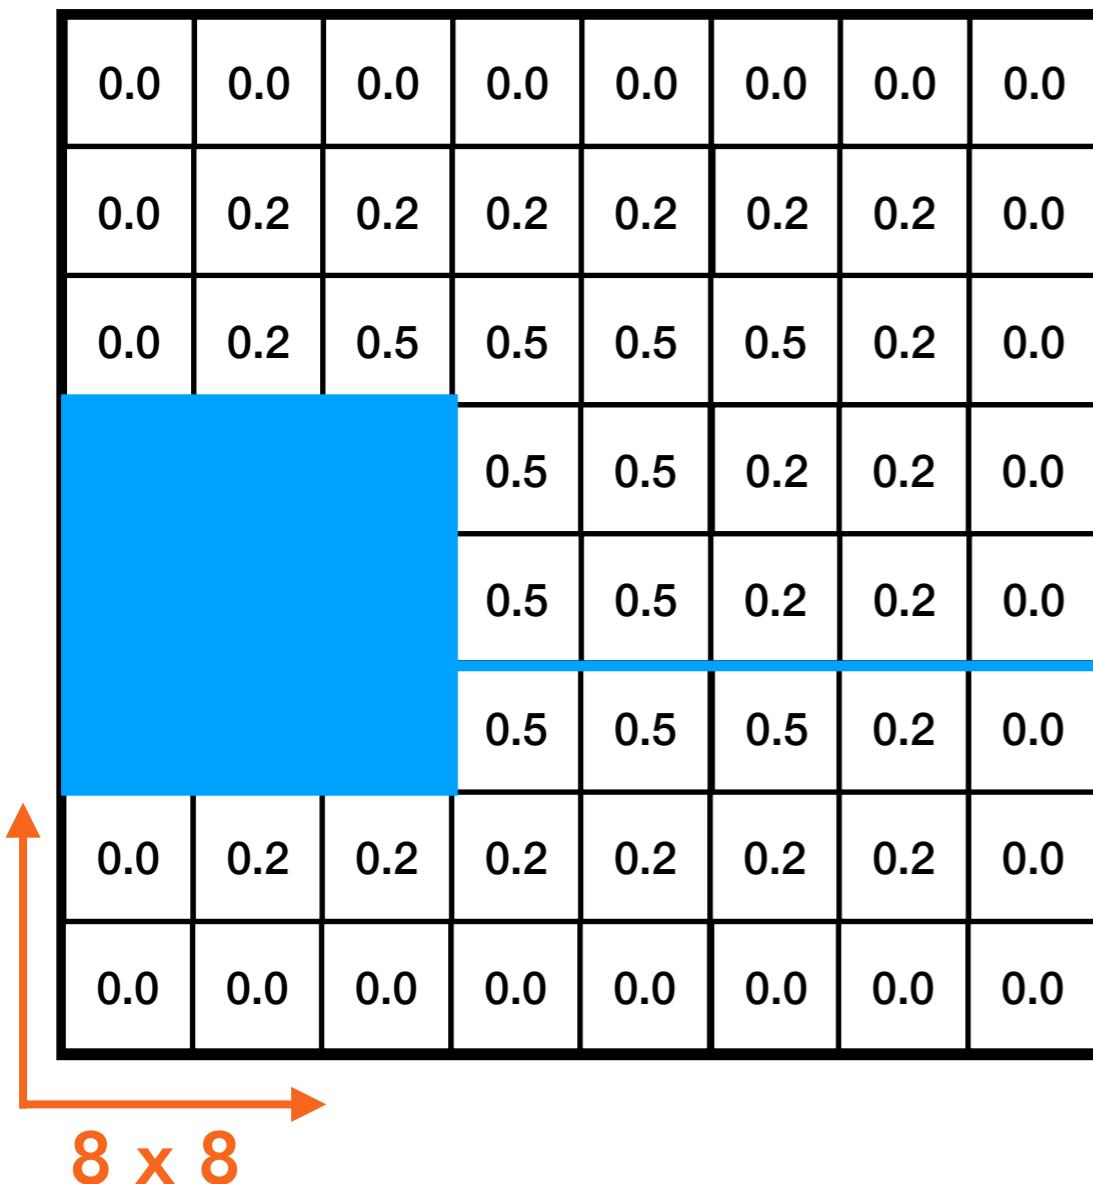
Multiple feature maps in each convolutional layer



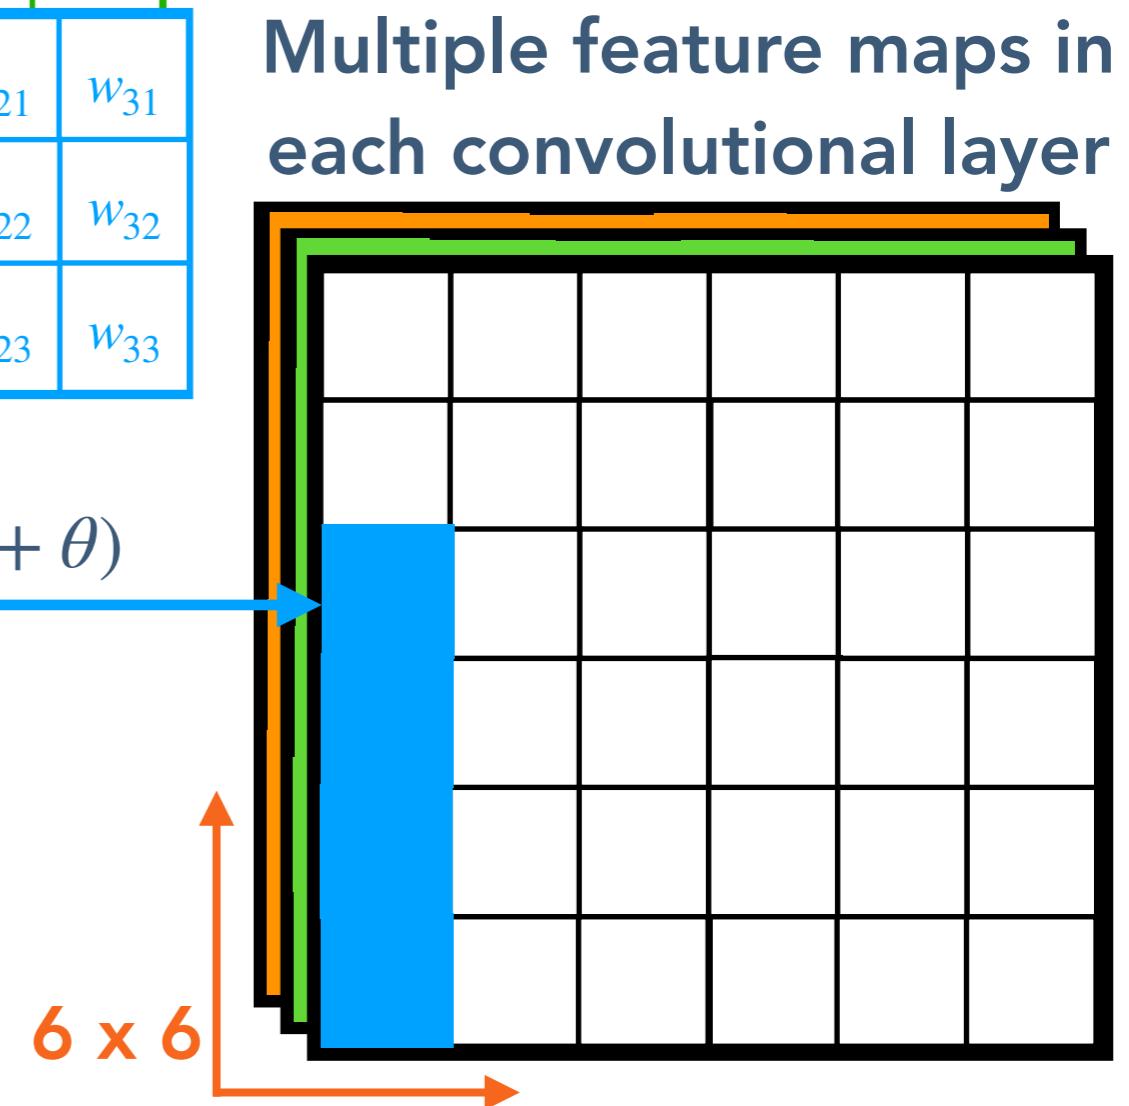


Dimension Reduction

- The feature map has reduced dimensionality compared to the input image -> loss of information.
- Comes from size of the filter and the stride used compared to the size of the input image.



$$f(W^T X + \theta)$$

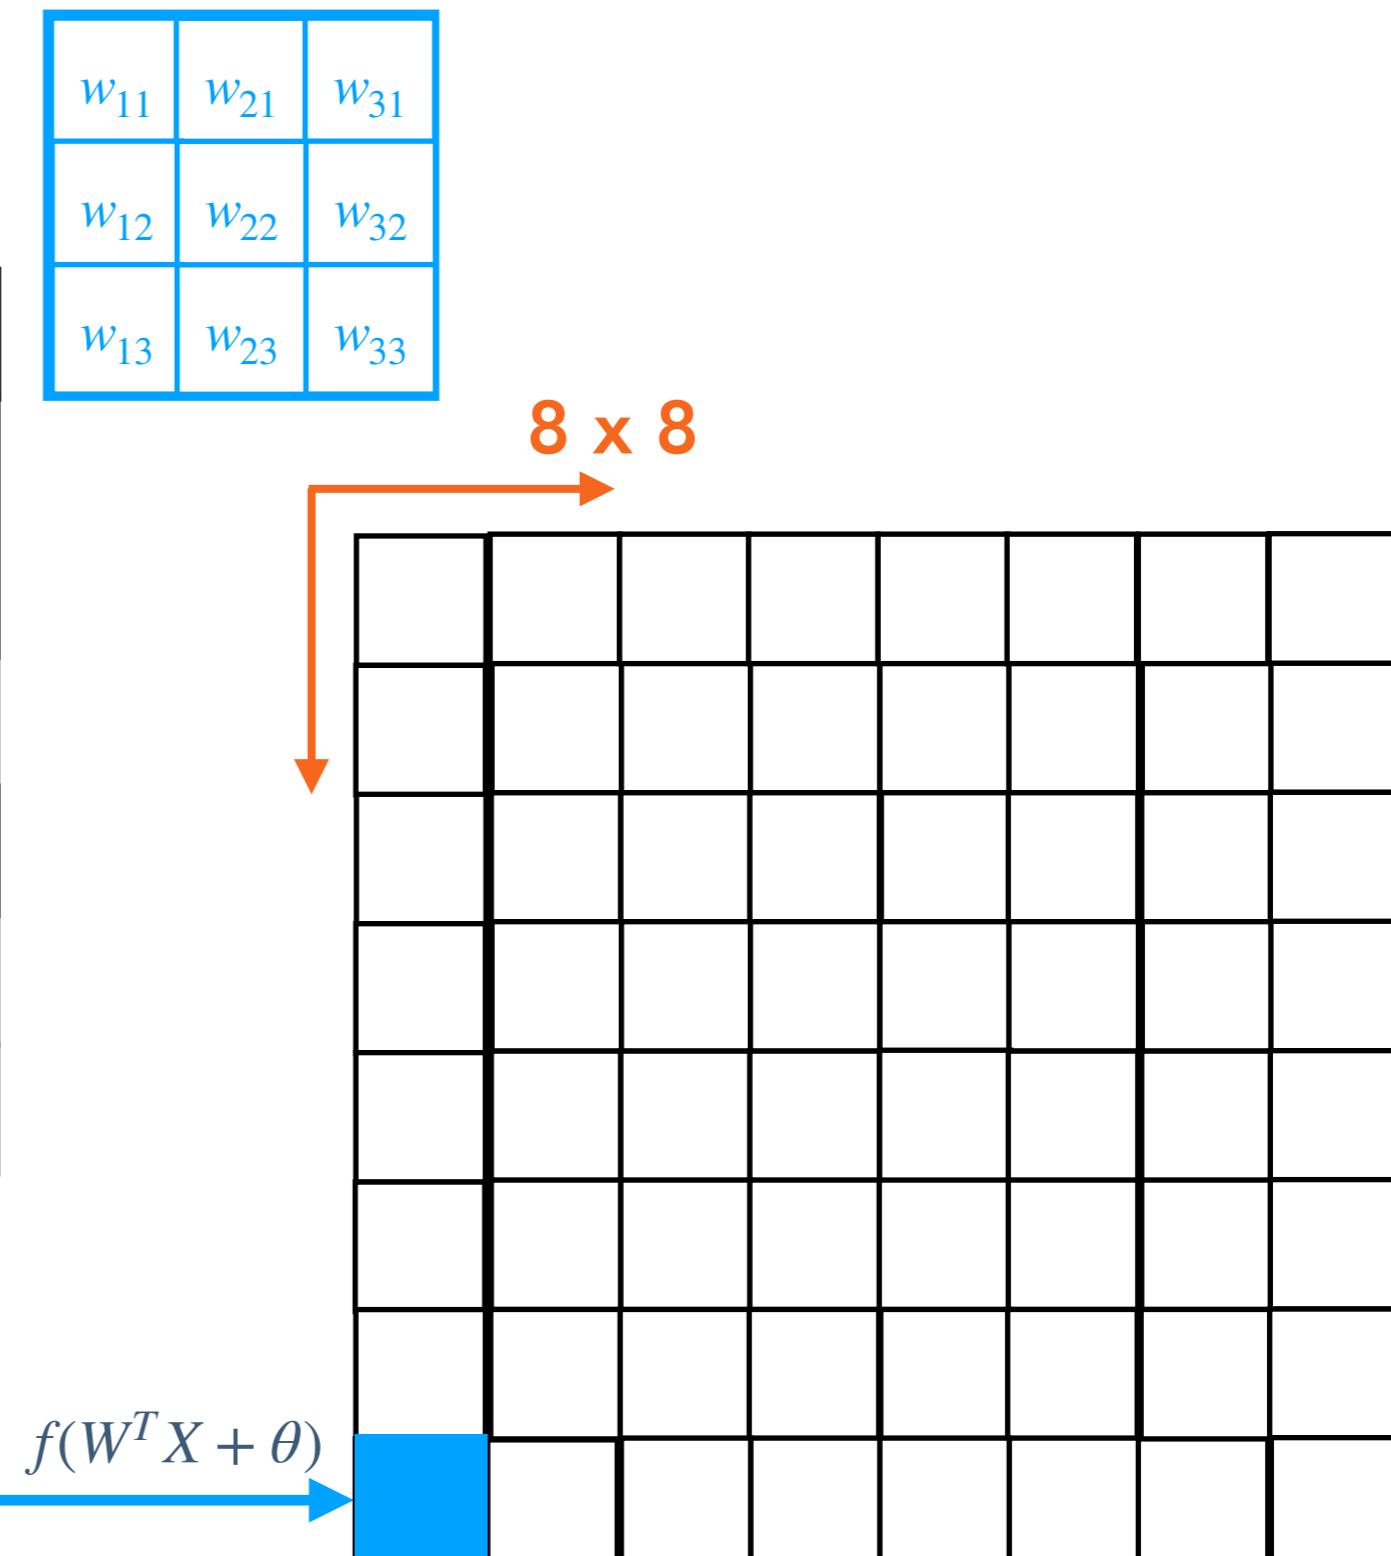




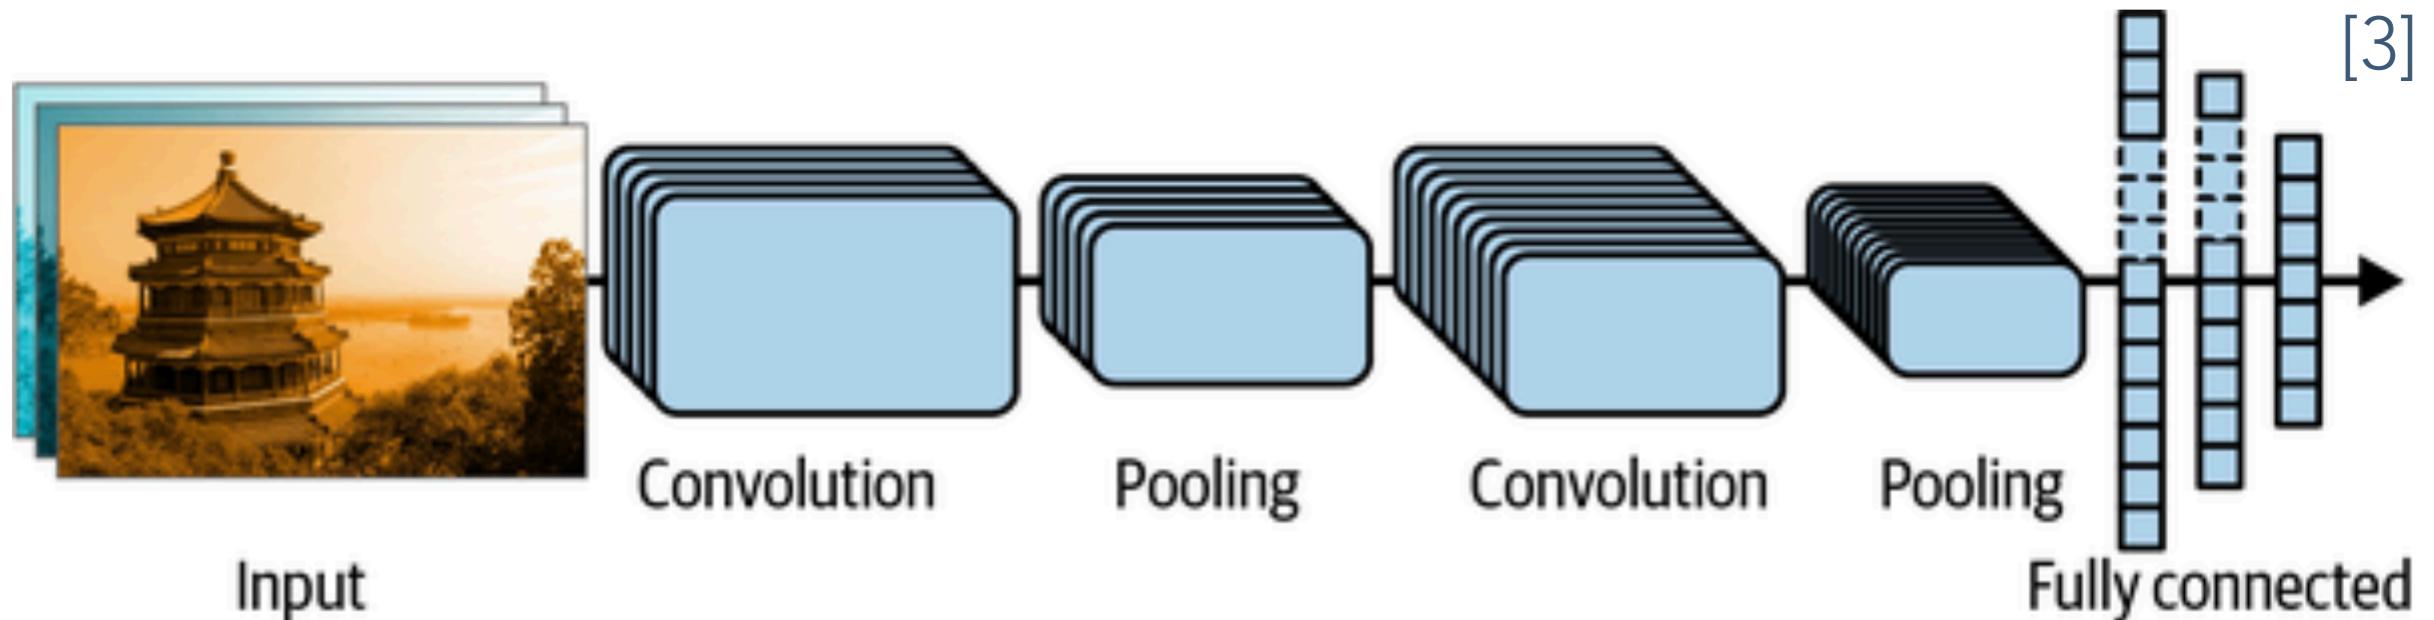
Dimension Reduction

- Loss of information is undesirable, use “zero padding” to recover original input size.
- No non-trivial information lost.

| | | | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.2 | 0.5 | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |



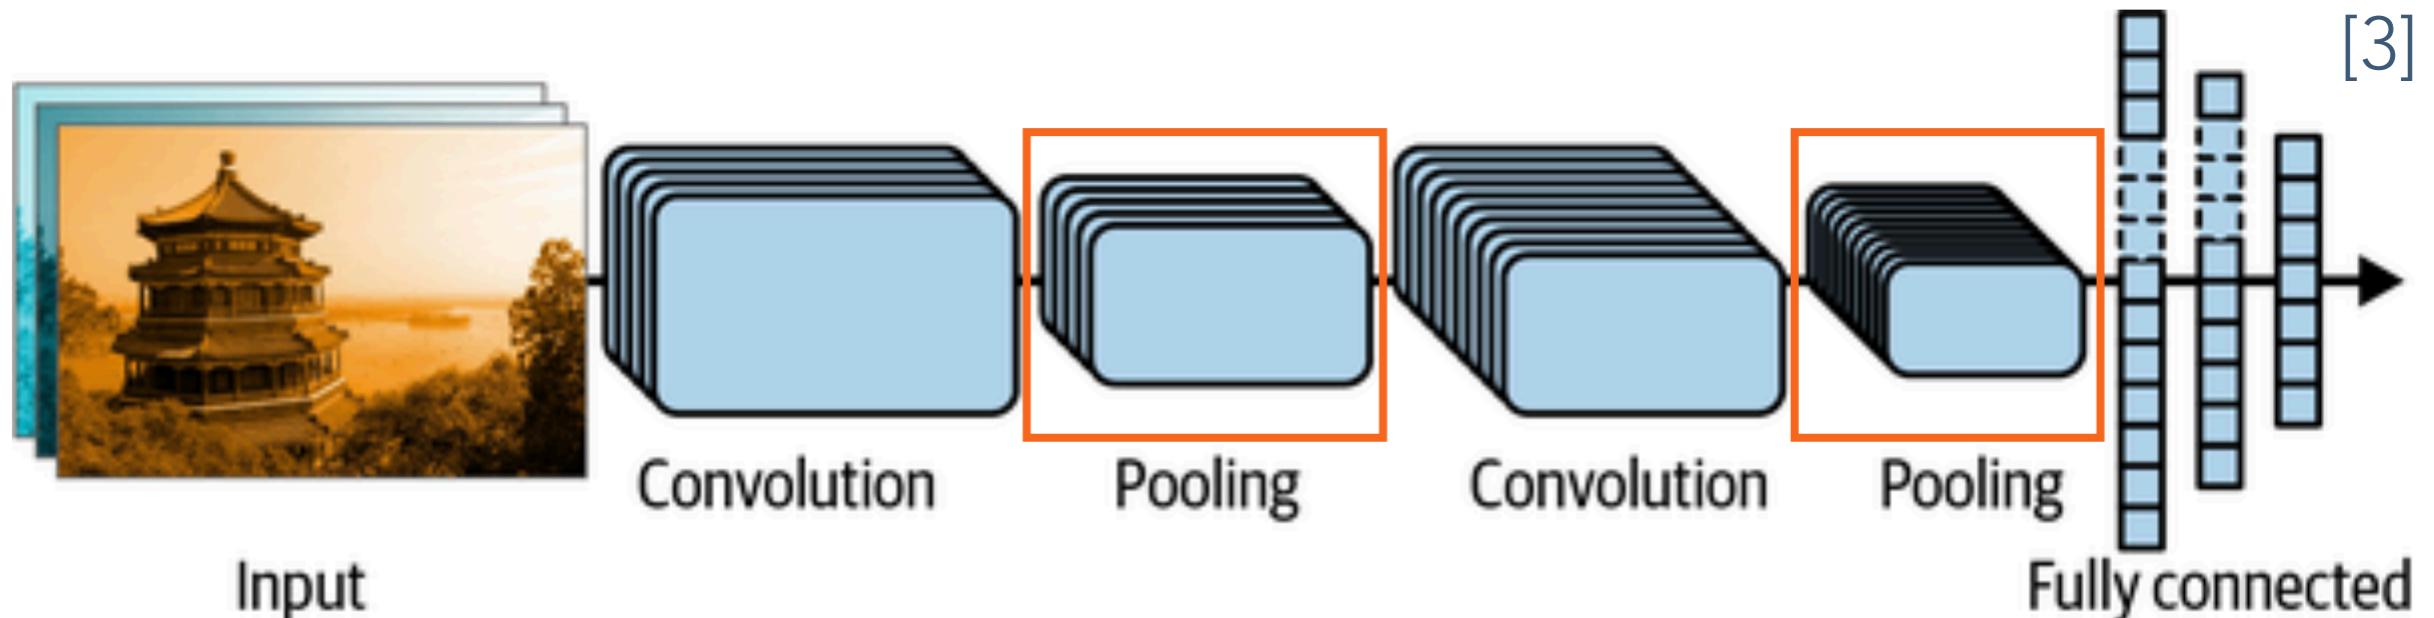
Typical CNN Architecture



- One or two convolution layers immediately after the input.
- Followed by a series of a pooling layer followed by a few convolutional layers.
- Finally standard, fully connected MLP to collate the convolution outputs as features and format the output for classification (via softmax).



Pooling



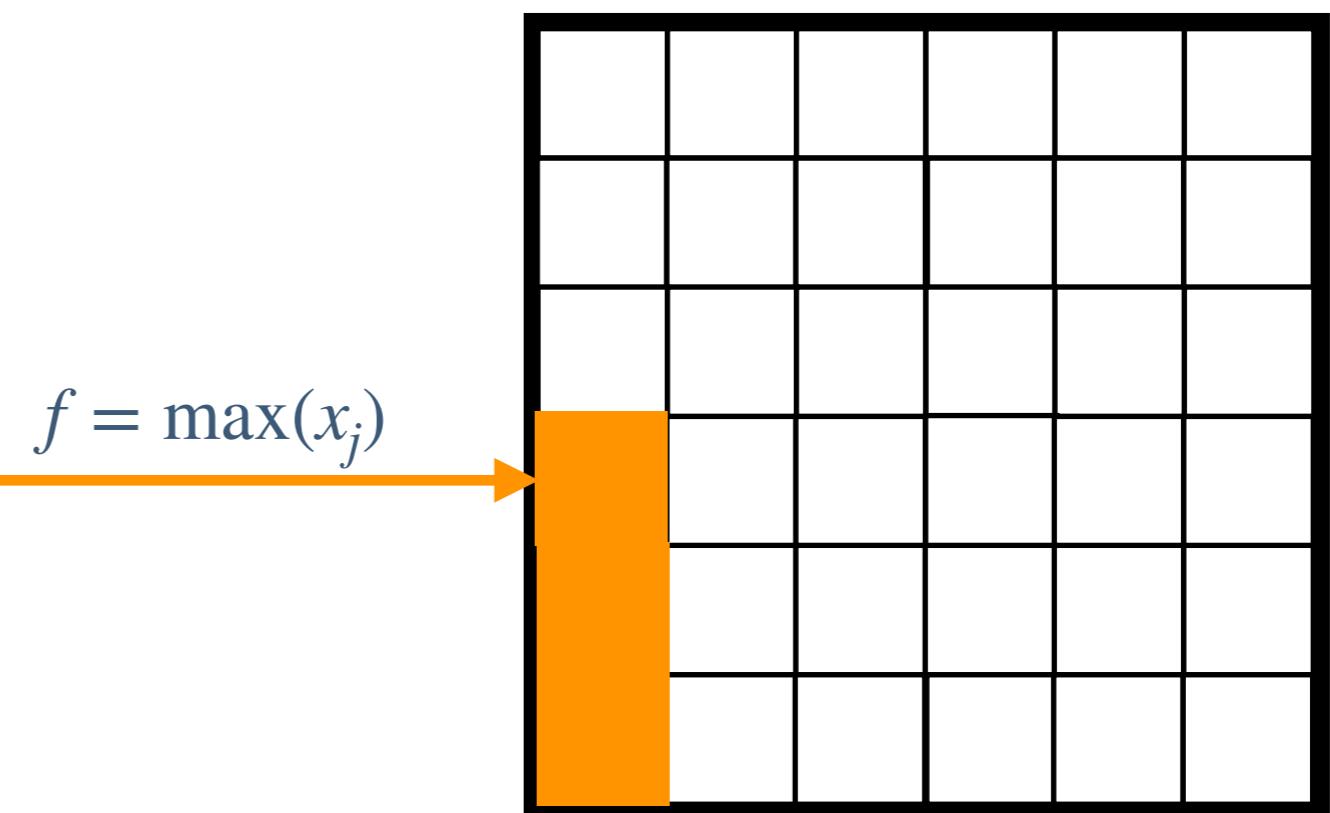
- Sometimes desirable to reduce the dimensionality of input image (or feature map):
 - ▶ Reduce the computational load.
 - ▶ Reduce memory usage.
 - ▶ Reduce the number of parameters (limiting the risk of overfitting).
- Often don't need the highest resolution possible to identify the shapes of objects -> use lower resolution representation and reach the same conclusion.
- This **subsampling** is the goal of the pooling layers.



Pooling Layers

- Define some filter (**kernel**) size, a stride and padding type then perform an operation on the pixels in the kernel's receptive field at each neuron to calculate:
 - ▶ their maximum value (**max pooling**).
 - Invariant under small changes - result will be the same post-pooling if input image is rotated or translated slightly -> **don't sweat the small things**.
 - Increase receptive field - summarising information in larger regions enables the network to capture more significant spacial patterns -> **see the bigger picture**.

| | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.5 | | | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | |
| 0.5 | | | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | |
| 0.5 | | | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 | |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 | |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |

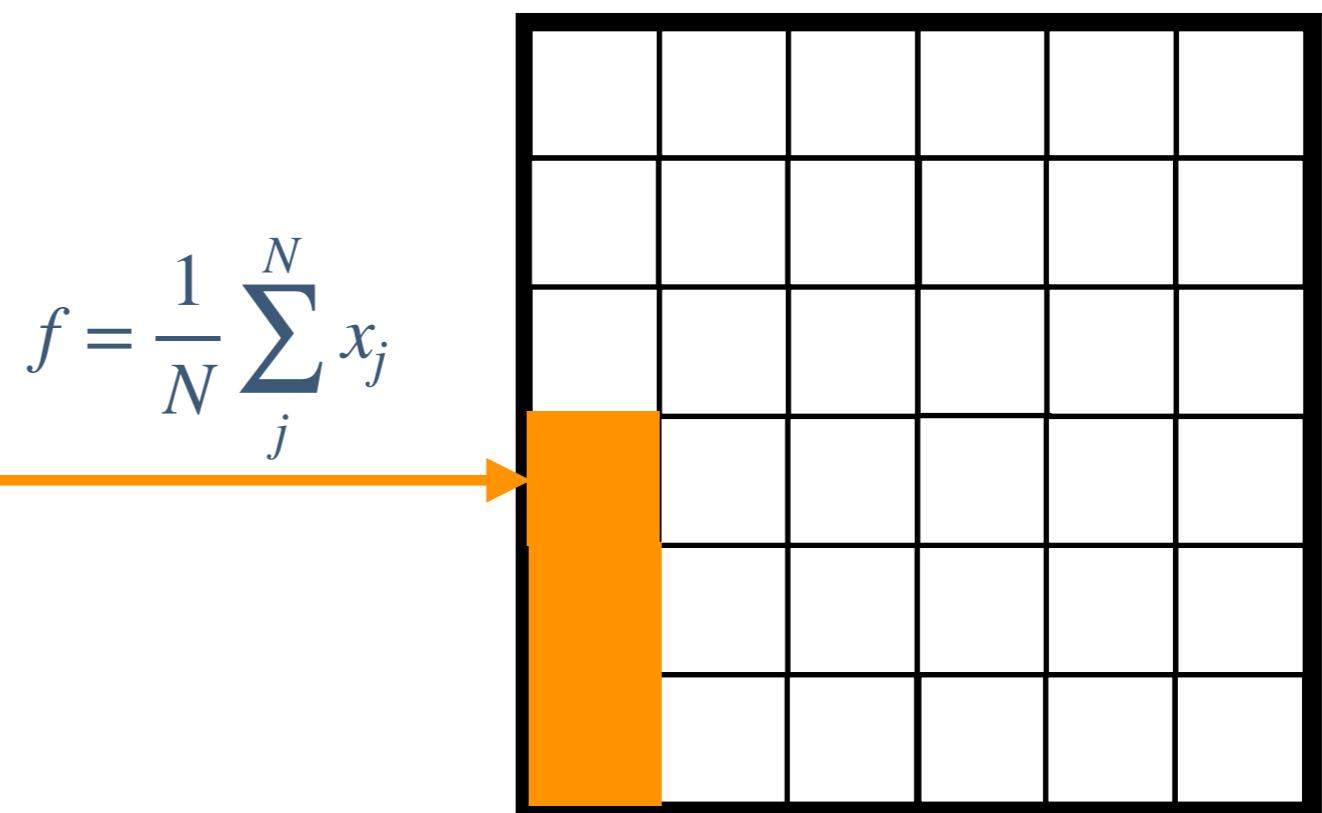




Pooling Layers

- Define some filter (**kernel**) size, a stride and padding type then perform an operation on the pixels in the kernel's receptive field at each neuron to calculate:
 - ▶ their mean value (**average pooling**).
 - Smoother feature representation - less variance in the pixels post-pooling.
 - Reduces sensitive to outliers -> regularisation effect - reducing overfitting by not letting any single pixel dominate.

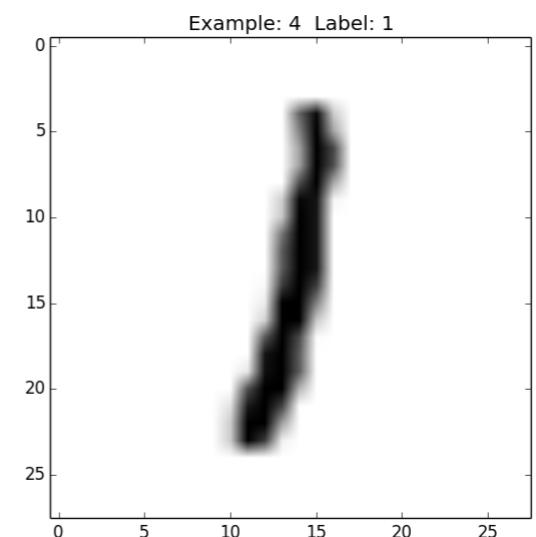
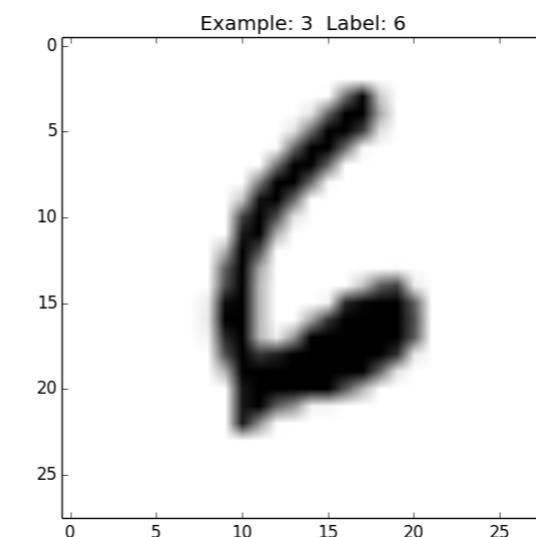
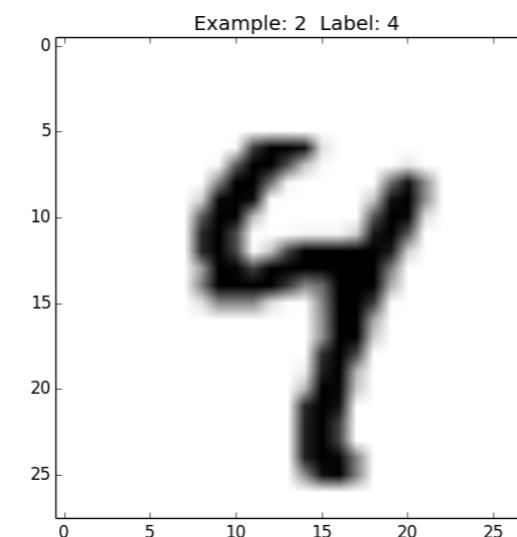
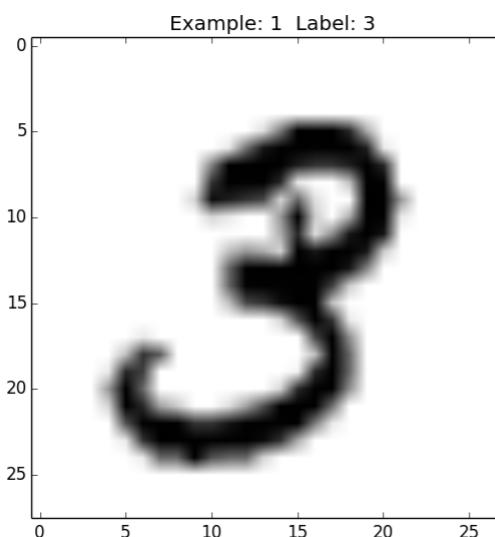
| | | | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|-----|-----|
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 |
| 0.0 | 0.2 | 0.5 | 0.5 | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 |
| 0.5 | | | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | |
| 0.5 | | | 0.5 | 0.5 | 0.2 | 0.2 | 0.0 | |
| 0.5 | | | 0.5 | 0.5 | 0.5 | 0.2 | 0.0 | |
| 0.0 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.2 | 0.0 | |
| 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | 0.0 | |





Today's Notebook

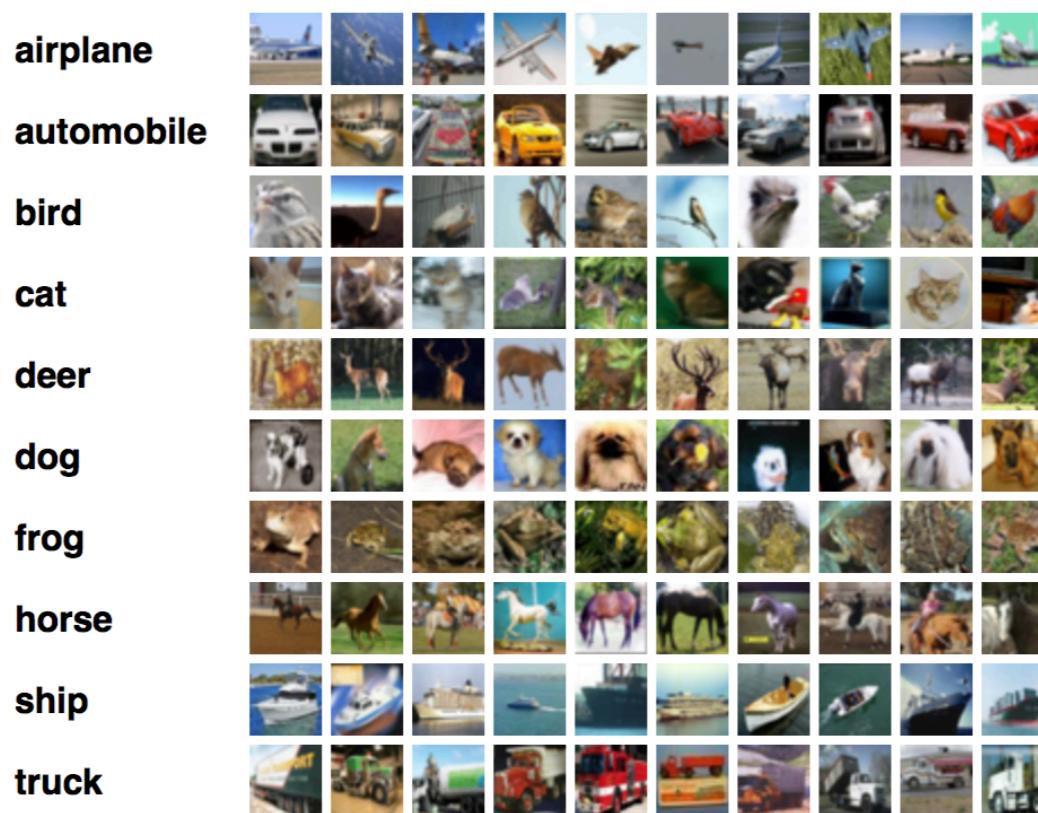
- Learned about all of the basic building blocks of a CNN.
- Discussed the process of convolution of images, why we do it and the need for padding.
- Introduced the concept of pooling, why it's useful and how to apply it.
- Put all of this together to build a CNN to classify images in two datasets.
 - ▶ **MNIST**: library of labelled handwritten greyscale numbers [2].





Today's Notebook

- Learned about all of the basic building blocks of a CNN.
- Discussed the process of convolution of images, why we do it and the need for padding.
- Introduced the concept of pooling, why it's useful and how to apply it.
- Put all of this together to build a CNN to classify images in two datasets.
 - ▶ **CFAR10:** library 60000 of labelled 32×32 pixel **colour** images in 10 classes [5].



**Trivial to extend CNN to images of arbitrary depth (colour = 3).

-> 3 fold increase in the number of weights in the network.

Accessing the Notebooks



Follow this link:

https://github.com/alexbooth92/workshop-UCDN_PracticalML.git

The screenshot shows a GitHub repository page for 'workshop-UCDN_PracticalML'. The repository has 12 commits, 1 branch, and 0 tags. The README file contains a brief description of the repository and a 'launch binder' button. The repository has 0 stars, 1 watching, and 0 forks. It also has 0 releases published.

A red arrow points to the 'launch binder' button in the README section.

About

A set of lectures and hands-on exercises introducing machine learning for a workshop At Universidad Católica del Norte, Chile.

Code

Commits

| Author | Message | Time | |
|-------------|--------------------------|---------------------------------------|----------------|
| alexbooth92 | Remove commit hash. | fc5f23b · 24 minutes ago | |
| | lecture_1 | Update train test to be better ratio. | 26 minutes ago |
| | lecture_2 | Initial commit of notebooks. | 2 weeks ago |
| | lecture_3 | Initial commit of notebooks. | 2 weeks ago |
| | lecture_4 | Initial commit of notebooks. | 2 weeks ago |
| | optional_python_revision | Binder test commit. | 2 weeks ago |
| | README.md | Update README.md | 2 weeks ago |
| | requirements.txt | Remove commit hash. | 24 minutes ago |

README

workshop-UCDN_PracticalML

A set of lectures and hands-on exercises introducing machine learning for a workshop At Universidad Católica del Norte, Chile. All material is based on a similar course put together by [Abbey Waldron](#).

You can run the code using TensorFlow 2.11 locally or online. To run, click on the following: [launch binder](#)

About

A set of lectures and hands-on exercises introducing machine learning for a workshop At Universidad Católica del Norte, Chile.

Code

Commits

| Author | Message | Time | |
|-------------|--------------------------|---------------------------------------|----------------|
| alexbooth92 | Remove commit hash. | fc5f23b · 24 minutes ago | |
| | lecture_1 | Update train test to be better ratio. | 26 minutes ago |
| | lecture_2 | Initial commit of notebooks. | 2 weeks ago |
| | lecture_3 | Initial commit of notebooks. | 2 weeks ago |
| | lecture_4 | Initial commit of notebooks. | 2 weeks ago |
| | optional_python_revision | Binder test commit. | 2 weeks ago |
| | README.md | Update README.md | 2 weeks ago |
| | requirements.txt | Remove commit hash. | 24 minutes ago |

Activity

0 stars

1 watching

0 forks

Releases

No releases published

[Create a new release](#)

Packages

No packages published

[Publish your first package](#)

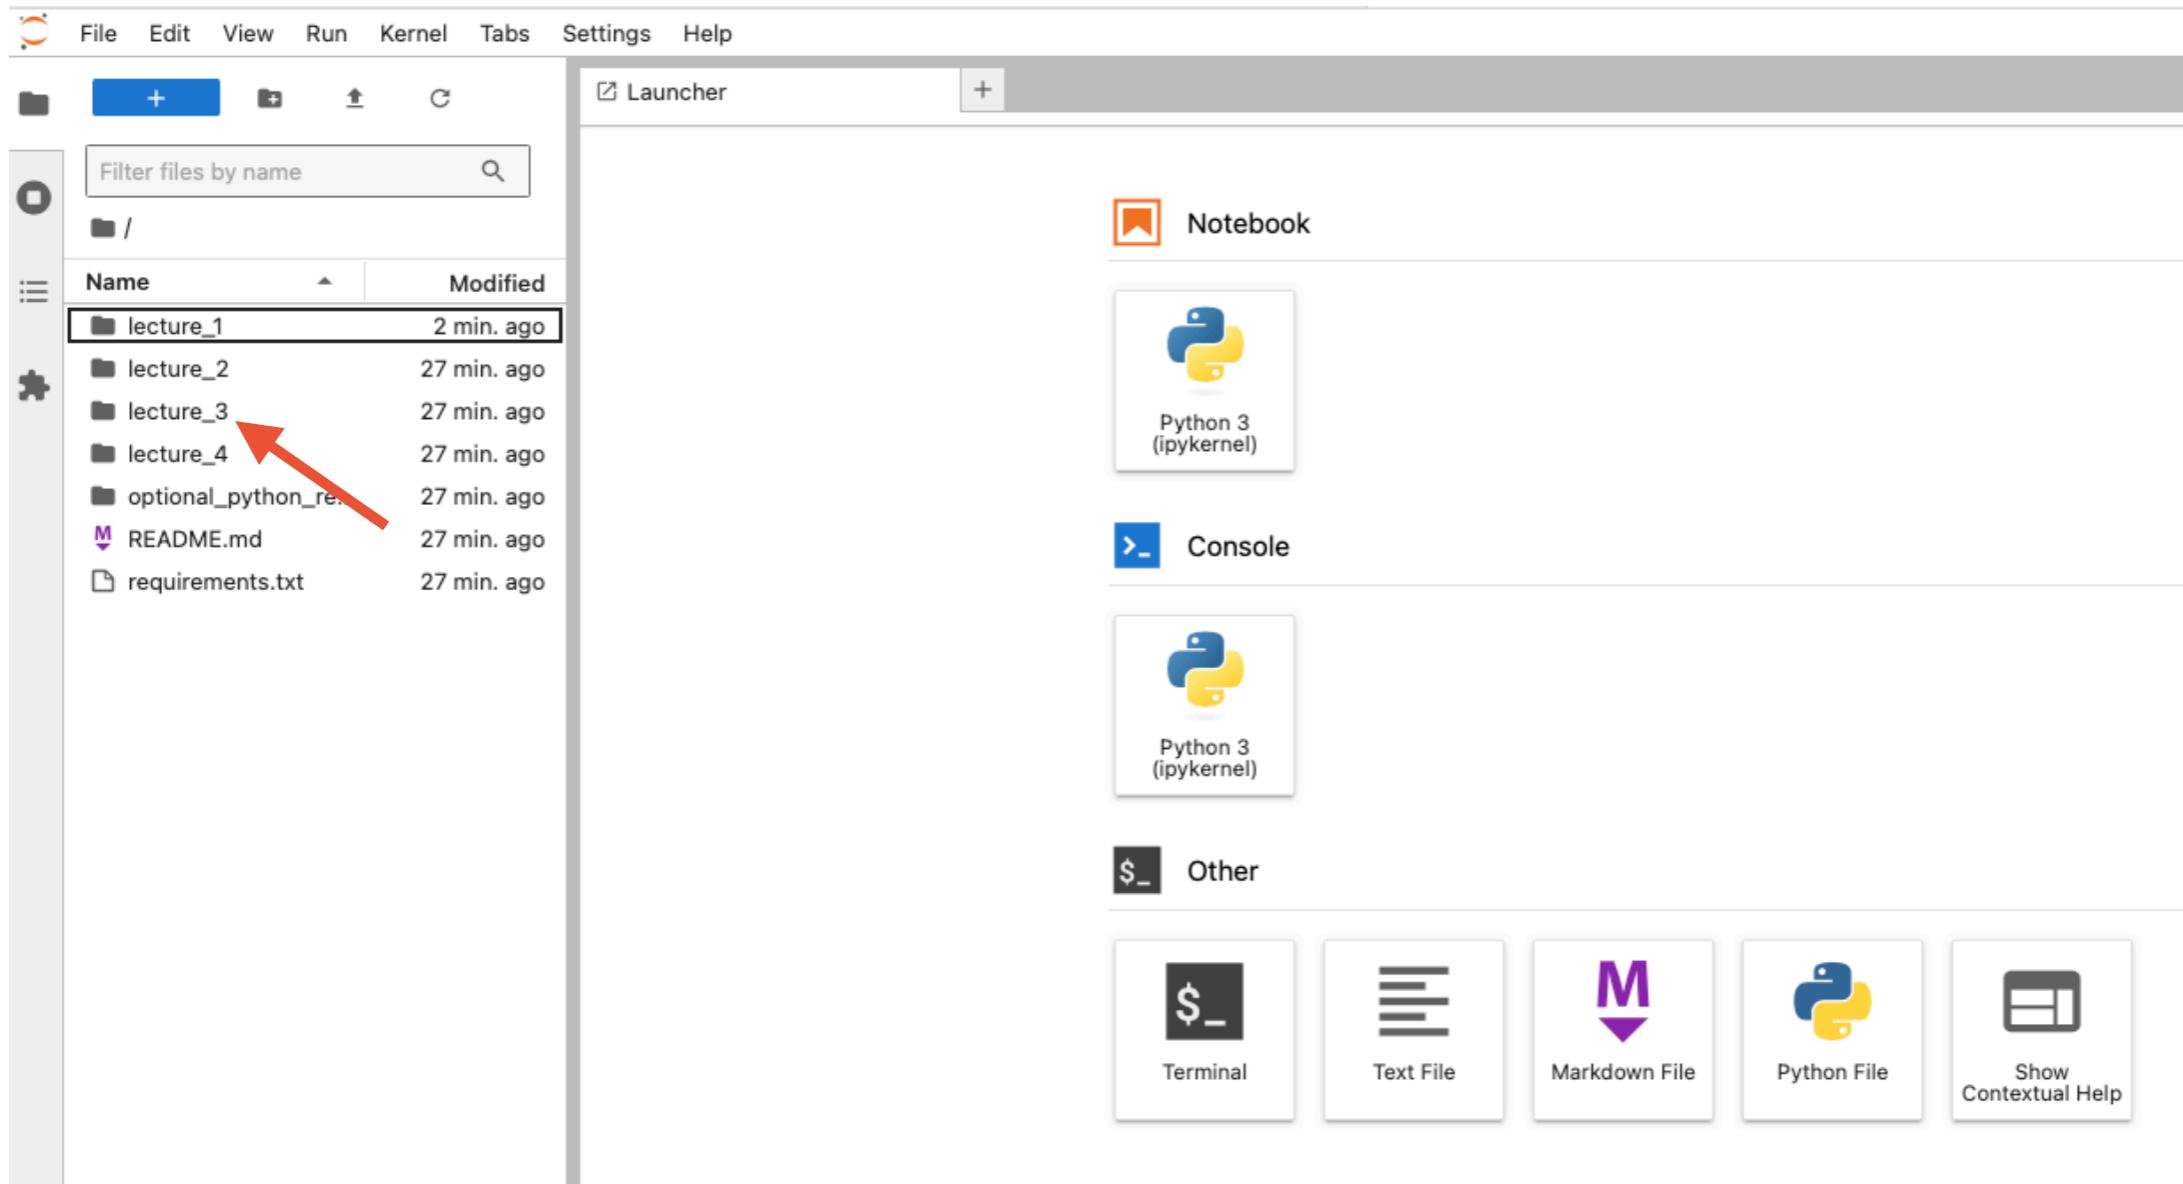
Languages

Jupyter Notebook 100.0%

Click here and be patient!



Accessing the Notebooks



Remember that dropout is your friend!

Have Fun!



References

- [1] <https://viso.ai/deep-learning/computer-vision-tasks/>
- [2] <http://yann.lecun.com/exdb/mnist/>
- [3] A. Géron. Hands on Machine Learning with Scikit-Learn & TensorFlow, 2017.
- [4] Neural Competition, Volume 22, Number 12, December 2010.
- [5] <https://www.cs.toronto.edu/~kriz/cifar.html>